

Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see <https://creativecommons.org/licenses/by-sa/2.0/legalcode>



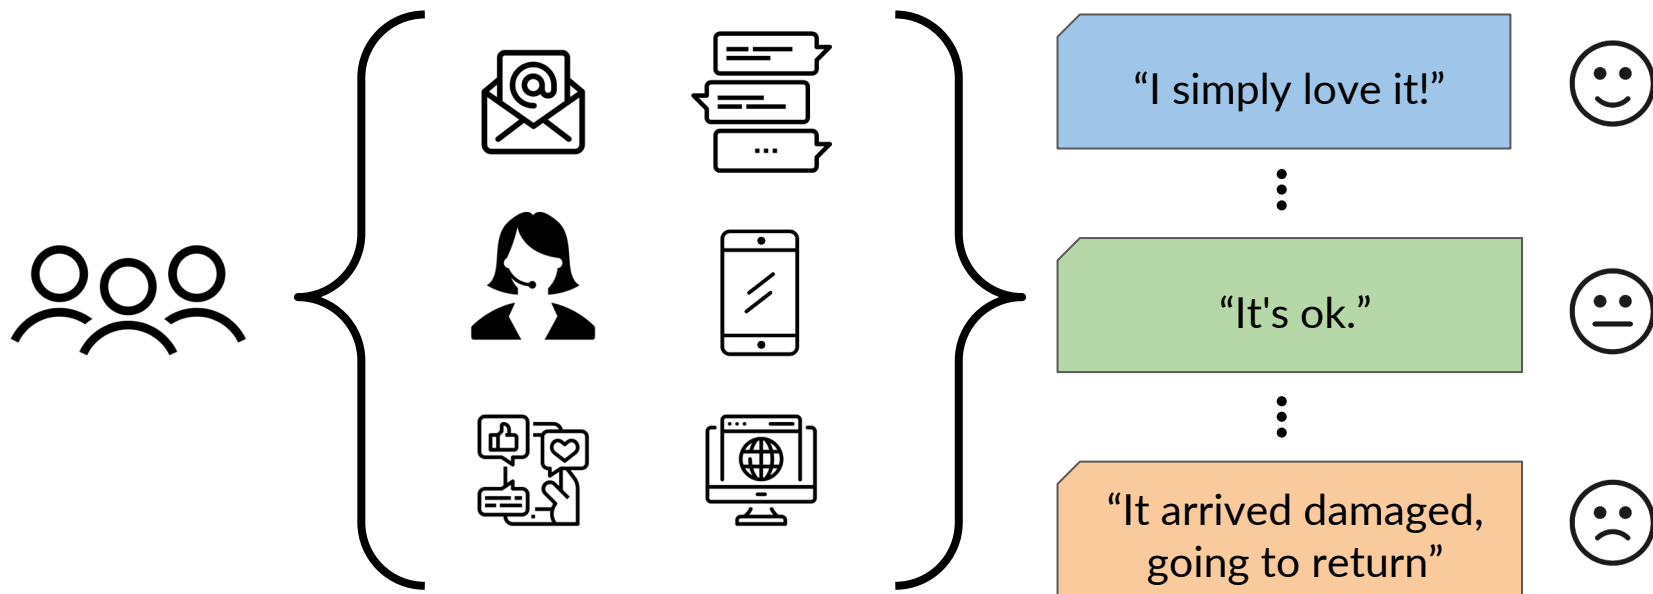
DeepLearning.AI



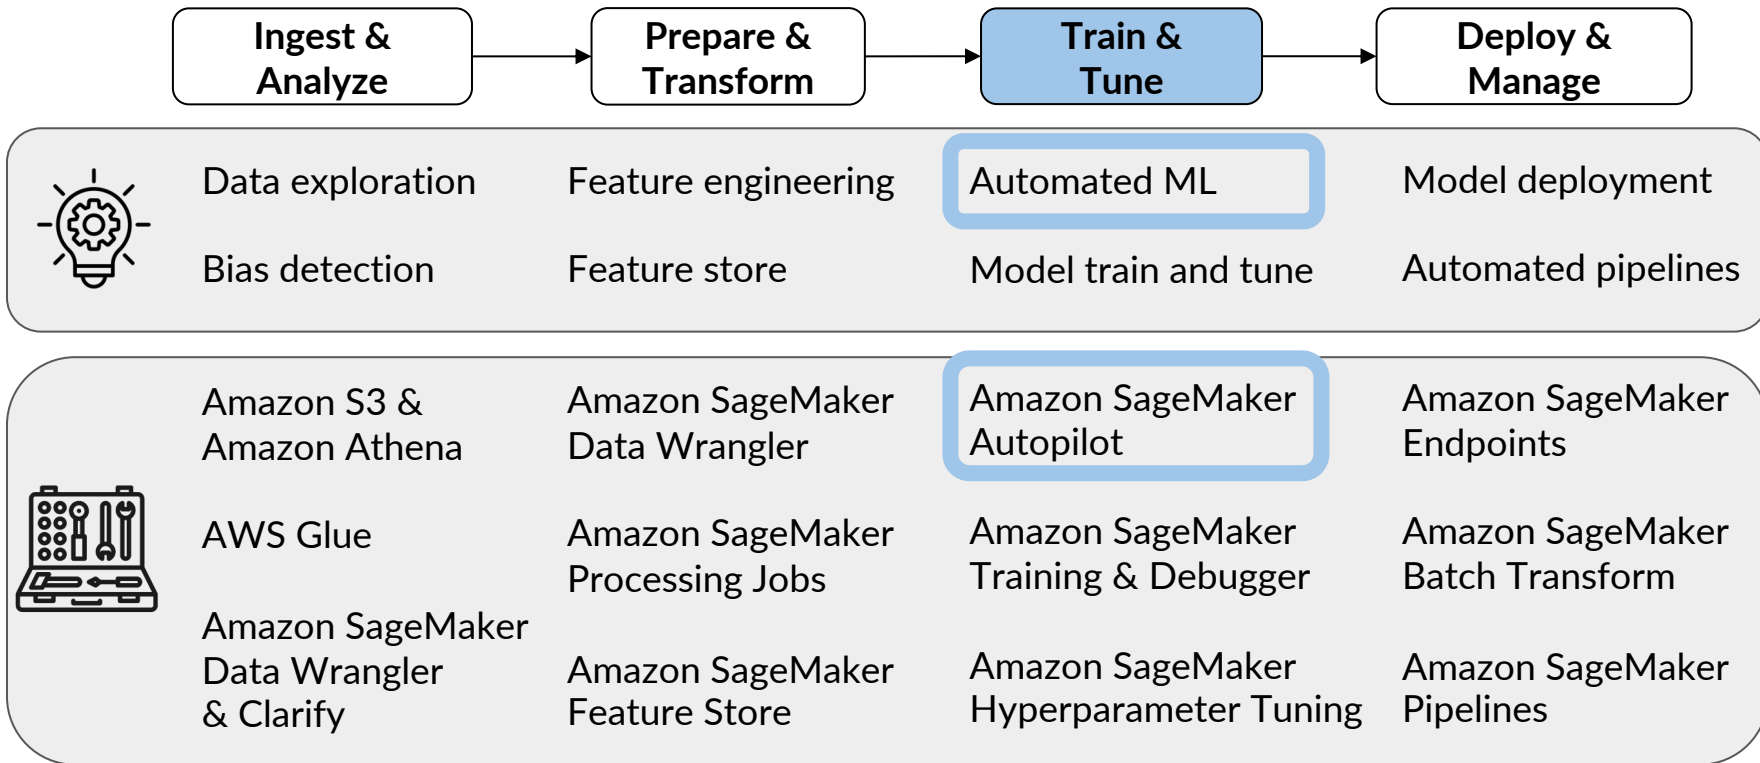
Practical Data Science

**Use AutoML to
Train a Text Classifier**

Multi-class classification for sentiment analysis of product reviews



Machine Learning Workflow





Automated Machine Learning

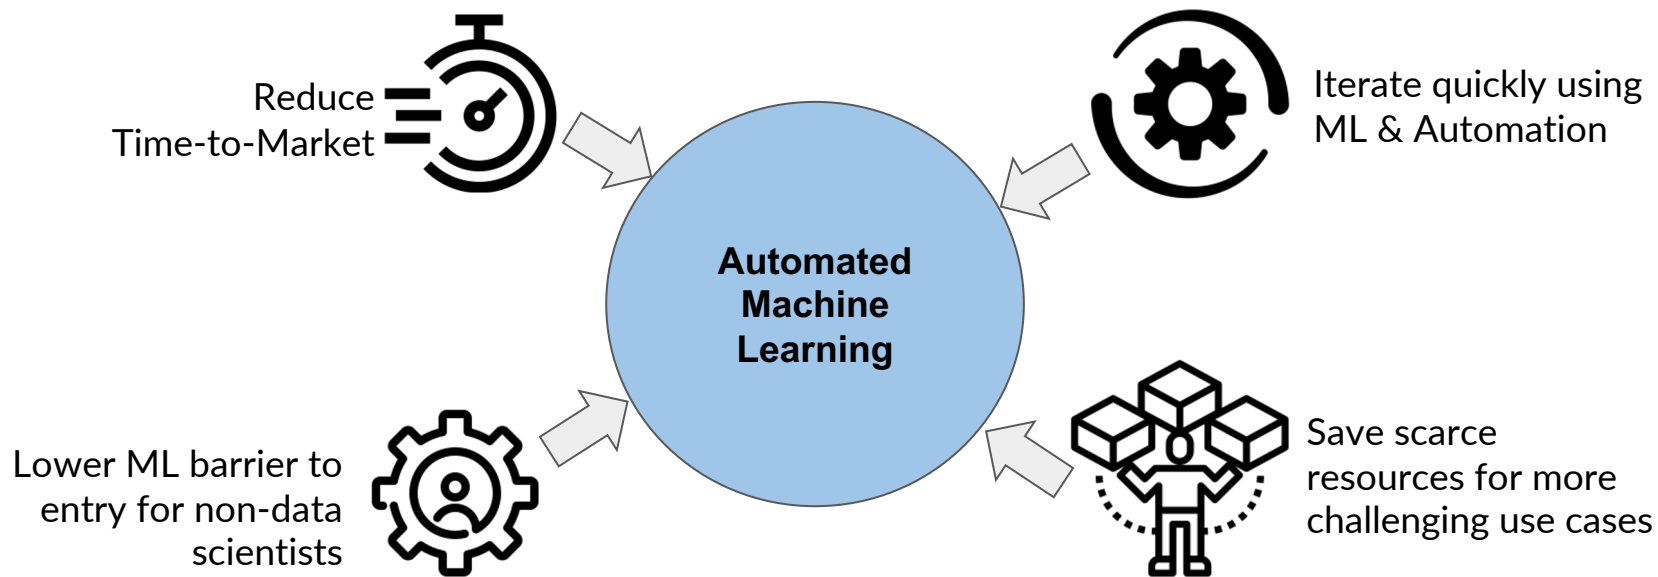
(AutoML)



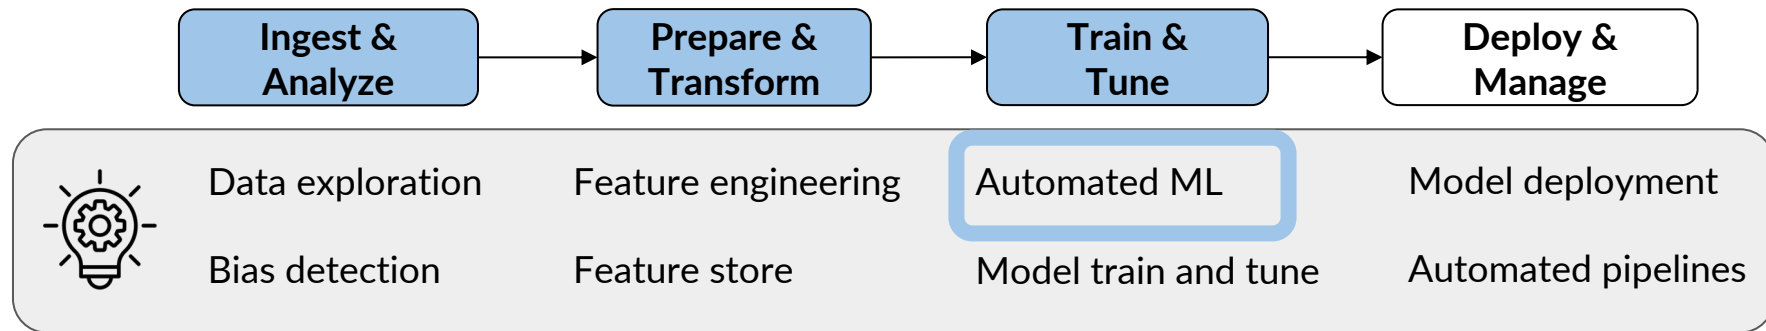
Model Building Challenges




Automated Machine Learning



Machine Learning Workflow

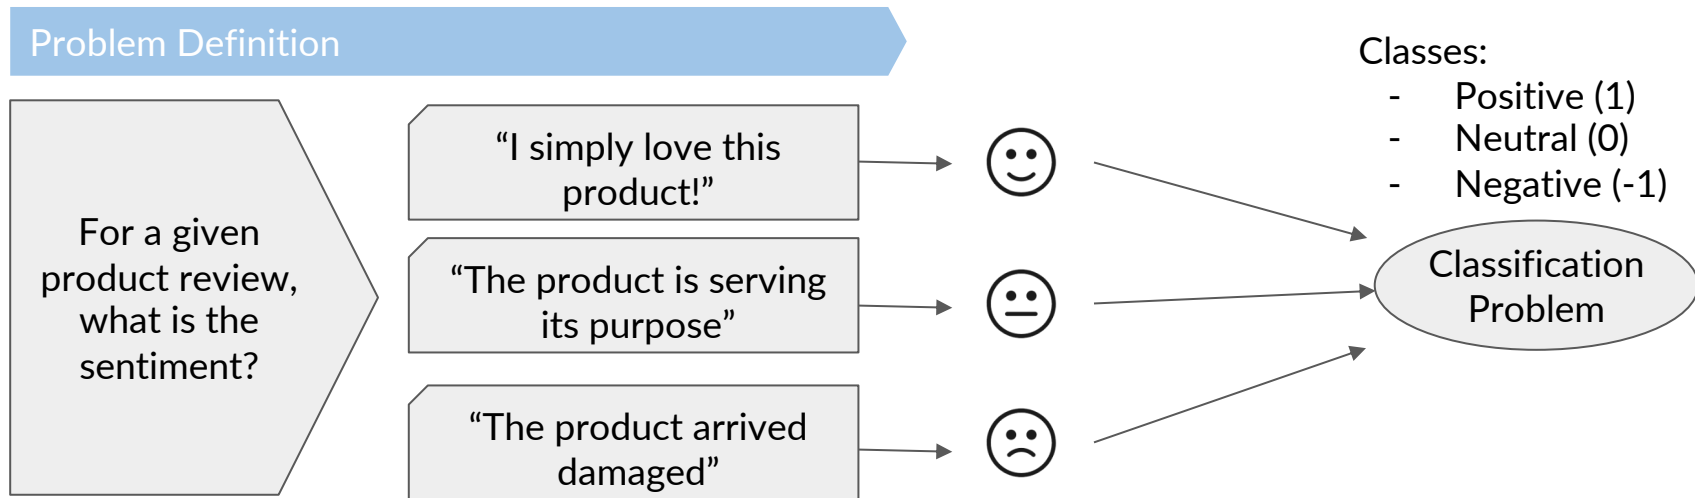


Data Preparation




Data Analysis:

Collecting statistics, such as missing entries, quantiles, skewness, correlation with the target.

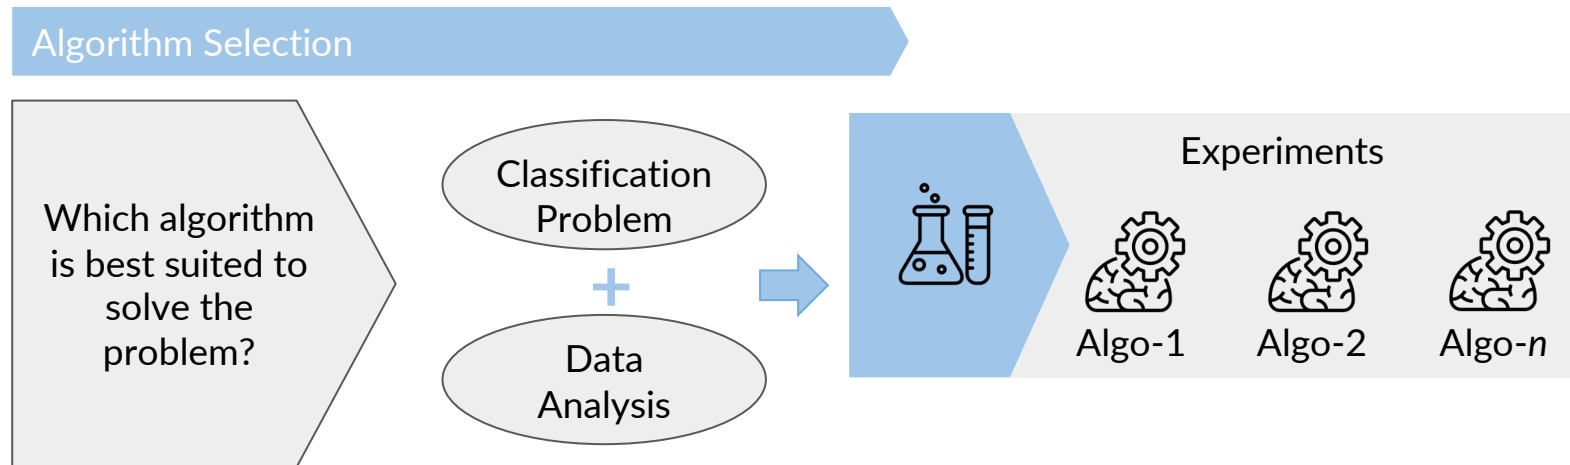


Data Preparation




Data Analysis:

Collecting statistics, such as missing entries, quantiles, skewness, correlation with the target.



Data Preparation




Data Analysis:

Collecting statistics, such as missing values, quantiles, skewness, correlation with the target.

Dataset Schema Detection

Numeric		Categorical	Numeric	
review_id		review_text	sentiment	
001		"I simply love this product!"	1	
002		"The product is serving its purpose"	0	
003		"The Product arrived damaged"	-1	

Data Preparation



Data Transformation:


How should data be transformed so that the model can predict as accurately as possible?

review_id	review_text	sentiment
001	"I simply love this product!"	1
002	"The product is serving its purpose"	0
003	"The Product arrived damaged"	-1





Too Many Unique Values
= Treat as Text

Data Preparation




Class Imbalance:

How to identify and handle potential class imbalance?




Data Preparation: Train and Validation Data Splits




Train-Validation Splits:

Splitting prepared data for model training, model performance, and final model evaluation

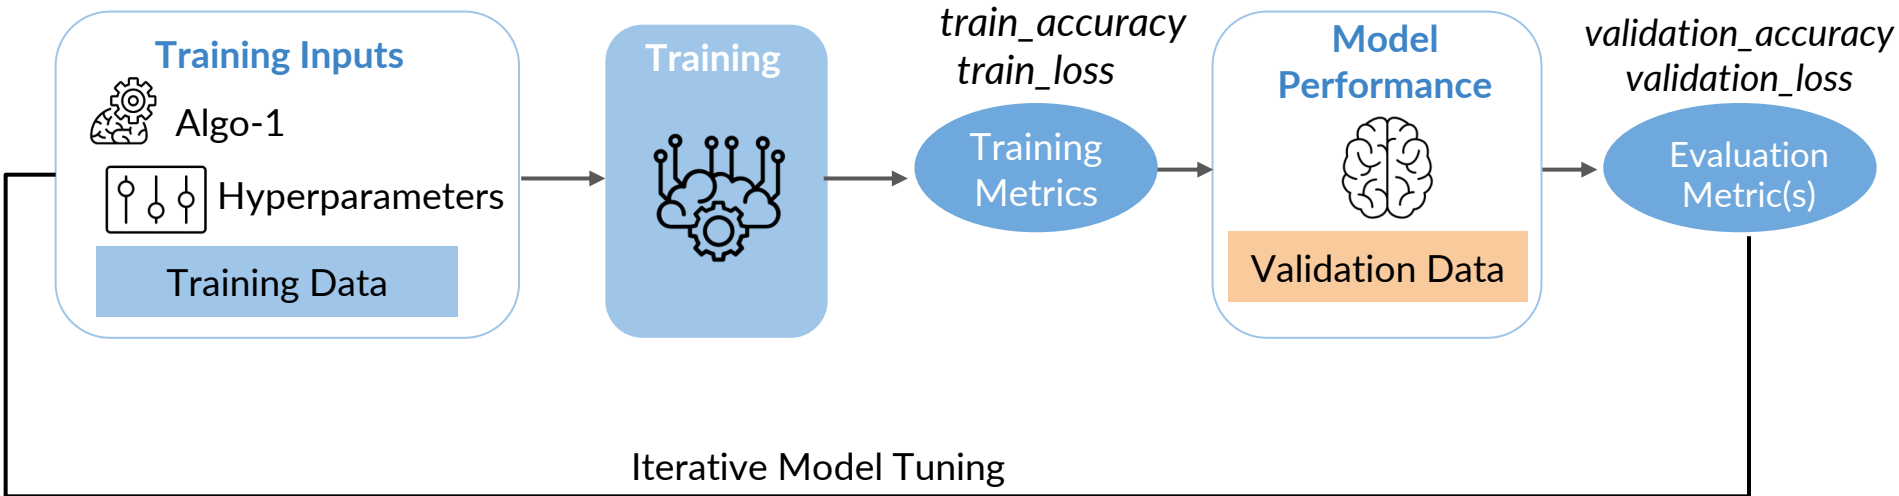


Model Training & Tuning




Model Training:

Fit the model to your data




AutoML

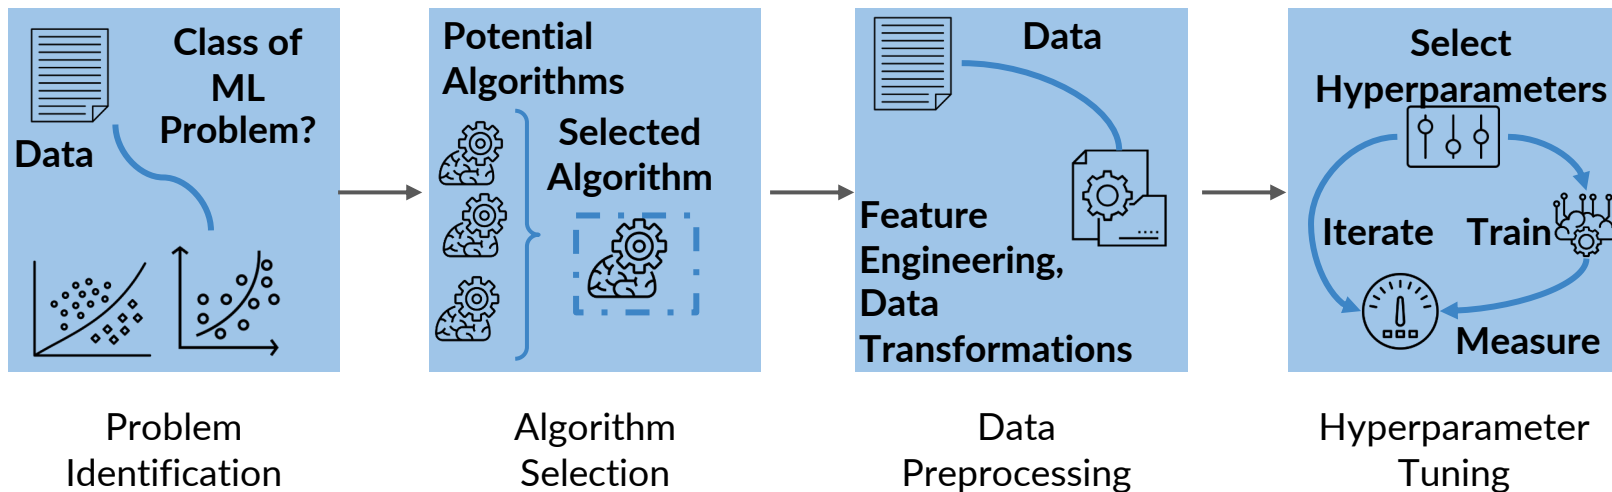


AutoML aims at automating the process of building a model

AutoML



AutoML aims at automating the process of building a model



Scenarios for AutoML

Build models without any ML expertise

- Empower more people in your organization: software developers, business people
- Let experts focus on **hard problems**

Experiment and build models at scale

- Thousands of data sets can be modeled without human intervention
- Let experts focus on **new problems**

Automate the majority of the work, then tweak

- Data cleaning, feature engineering, feature selection, etc.
- Let experts focus on high value tasks such as **domain knowledge**, and **error analysis**.

Transparency and Control are Important

Get the **best model** only

- Hard to understand it
- Hard to reproduce it manually

Get the **best model, all candidates, full source code**

- Understand how the model was built
- Keep tweaking for extra performance



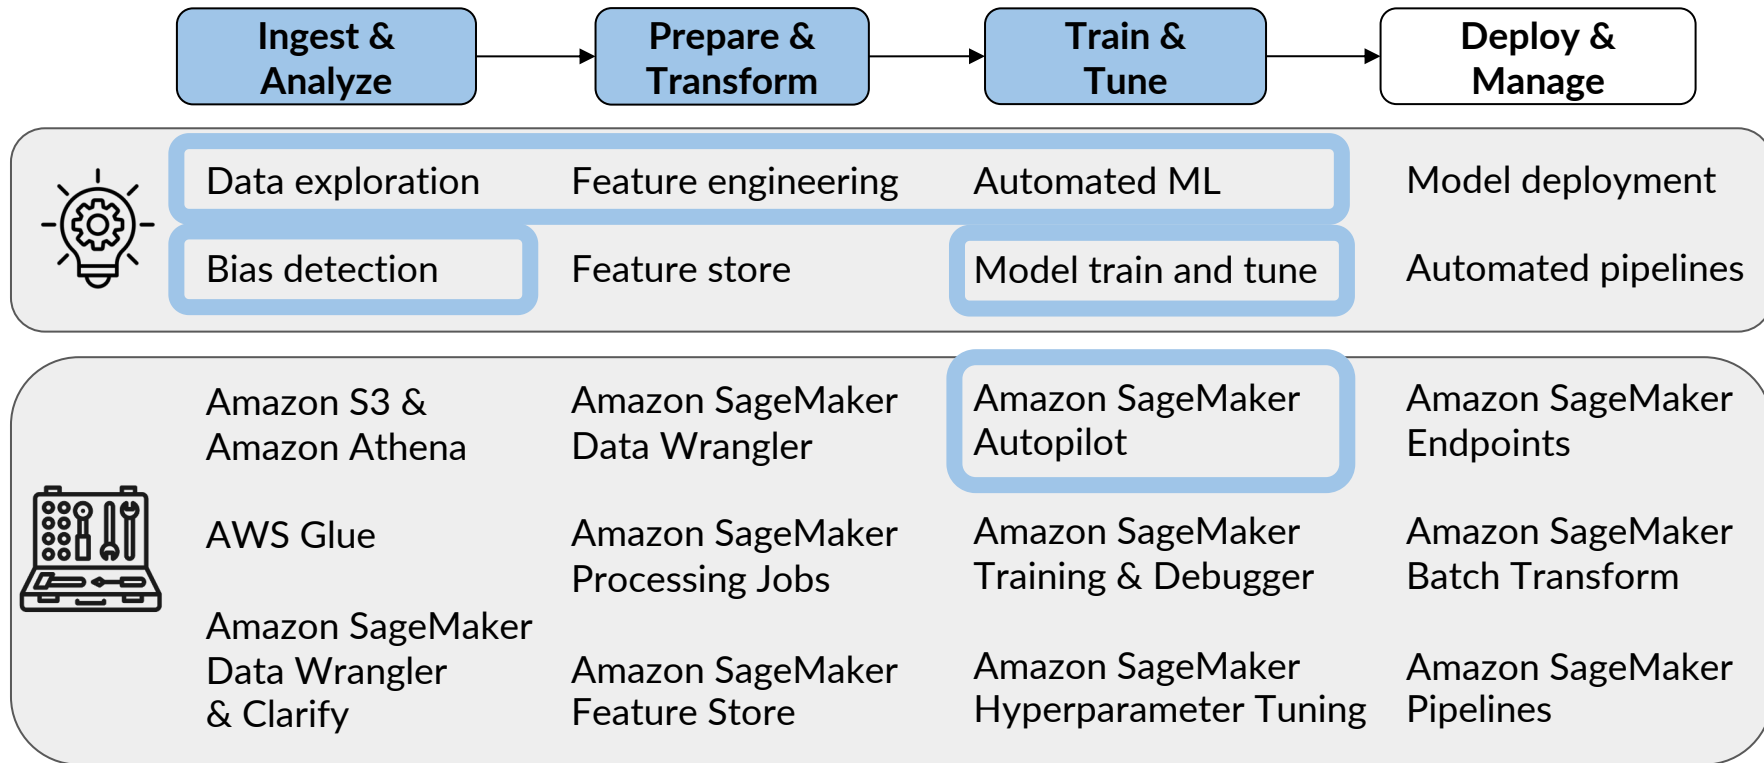
AutoML

with
Amazon SageMaker Autopilot

Introduction

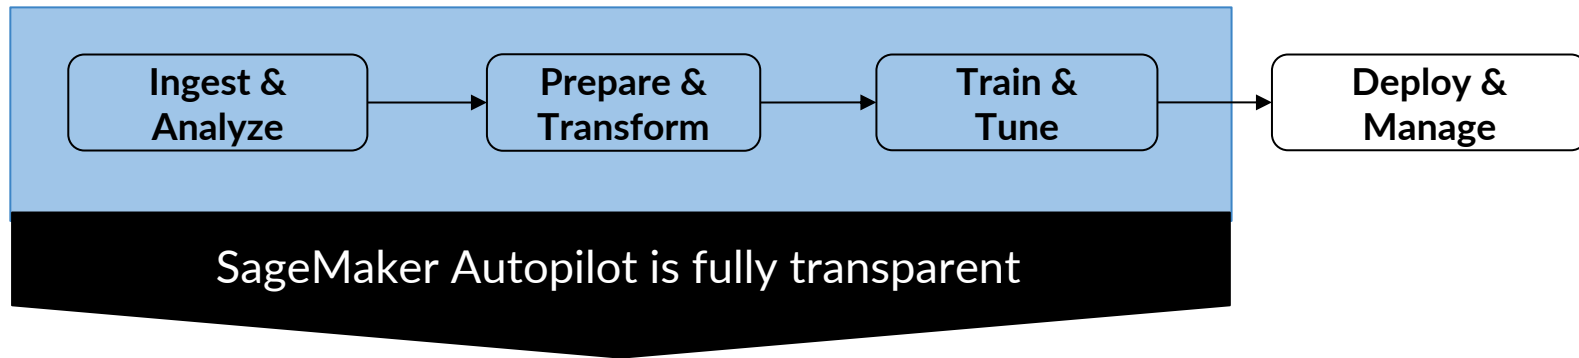


Machine Learning Workflow



AutoML with Amazon SageMaker Autopilot

Amazon SageMaker Autopilot covers all steps:



Automatically
Generated



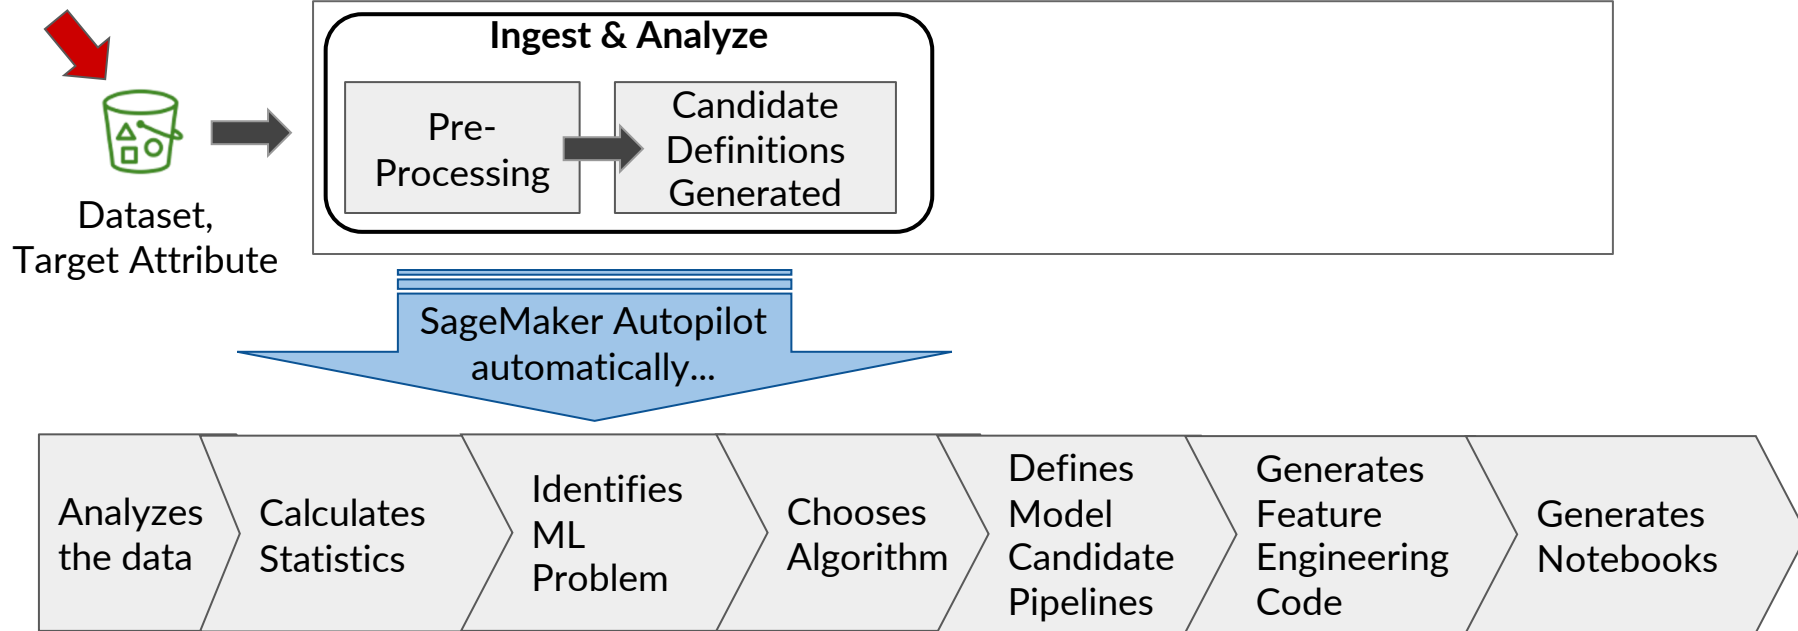
Notebooks



Code

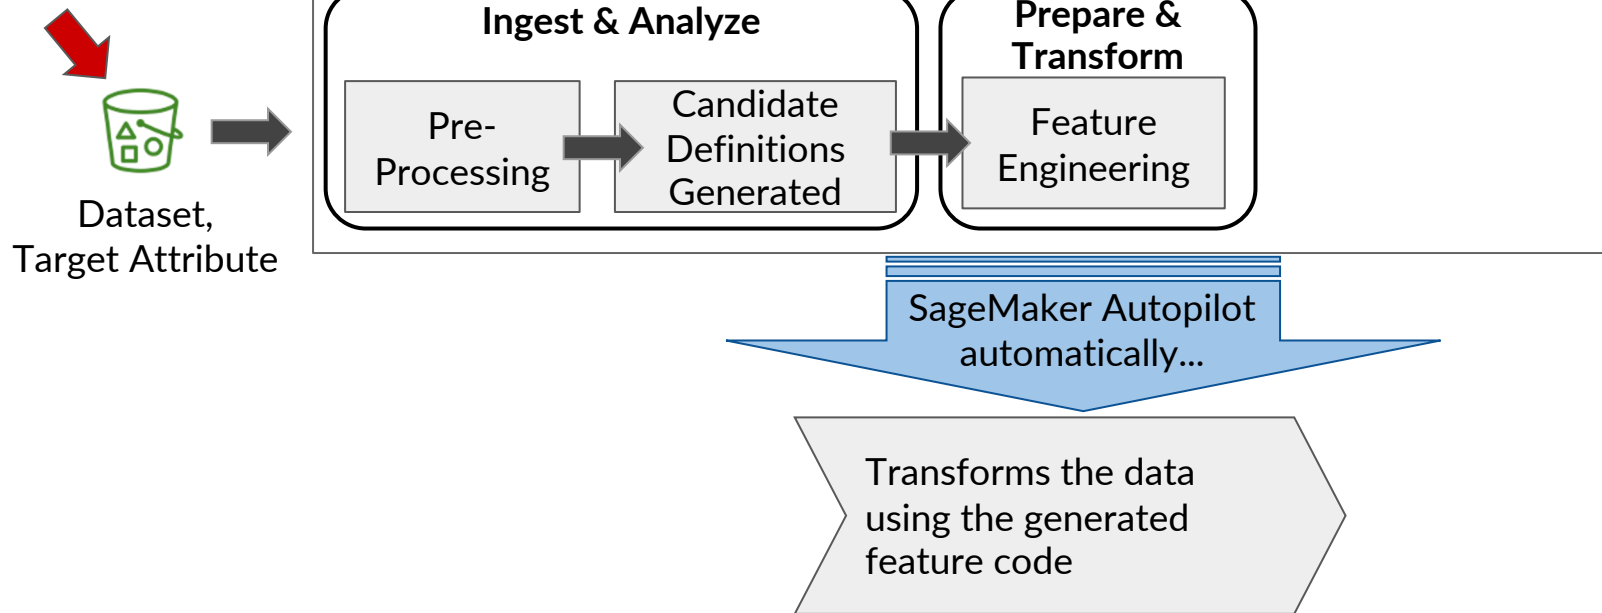
Amazon SageMaker Autopilot at a High-Level

Share your tabular dataset
in a S3 bucket



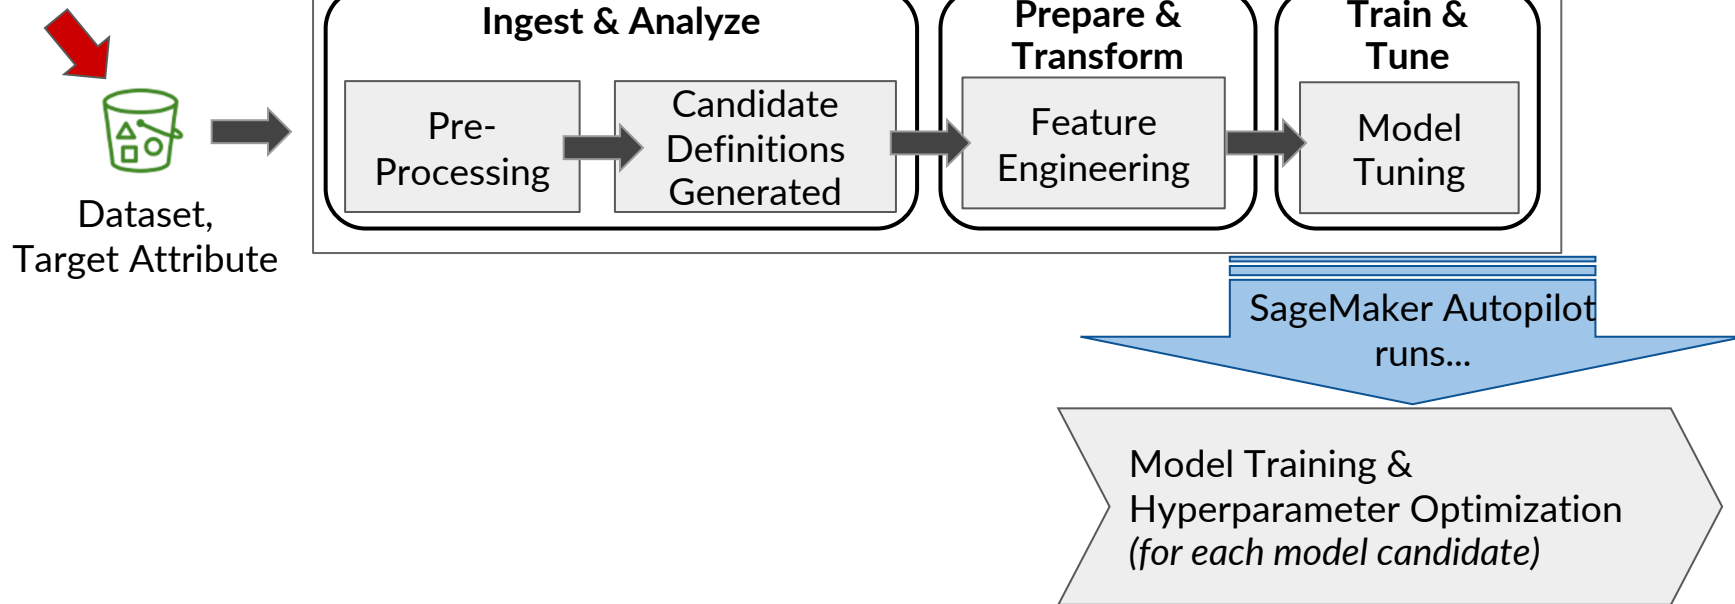
Amazon SageMaker Autopilot at a High-Level

Share your tabular dataset in a S3 bucket




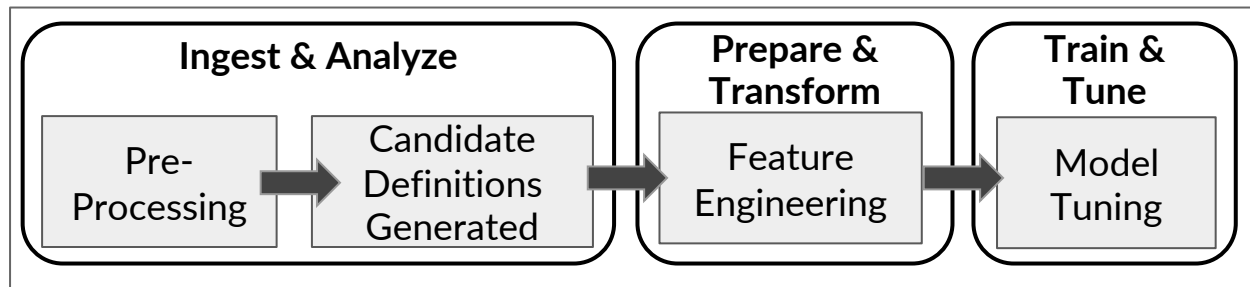
Amazon SageMaker Autopilot at a High-Level

Share your tabular dataset
in a S3 bucket



Amazon SageMaker Autopilot at a High-Level

Share your tabular dataset
in a S3 bucket



SageMaker Autopilot
shares...

- All metrics
- Leaderboard of model candidates
- Notebooks
- Code

AutoML

with
Amazon SageMaker Autopilot

Running Experiments



Amazon SageMaker Autopilot Notebook Overview

Use Case: Analyze Customer Sentiment

Goal: Use SageMaker Autopilot to find the optimal feature transformations, algorithm, and hyperparameters to produce a best performing model allowing us to predict our **label (sentiment)** based on **product reviews (review_body)**

sentiment	review_body
-1	This is bad.
0	This is OK.
1	This is great!

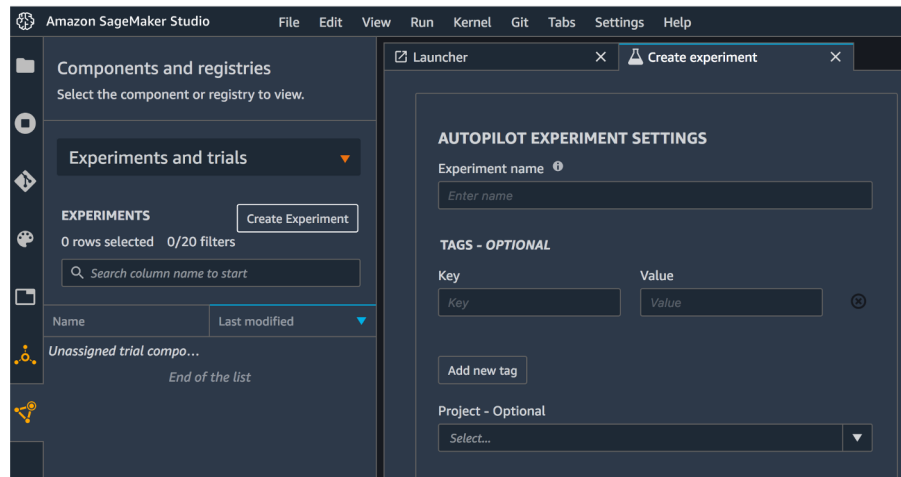
Interacting with Amazon SageMaker Autopilot



~OR~

Programmatically:

1. AWS CLI
2. AWS SDK
3. Amazon SageMaker Python SDK



Amazon SageMaker Studio

Launch the Amazon SageMaker Autopilot Job

```
automl = sagemaker.automl.automl.AutoML(  
    target_attribute_name=...,  
    output_path=...,  
    max_candidates=3,  
    role=role,  
    max_runtime_per_training_job_in_seconds=1200,  
    total_job_runtime_in_seconds=7200 # max automl job runtime in seconds  
)  
  
automl.fit(  
    inputs=...,  
)
```

Attribute to predict

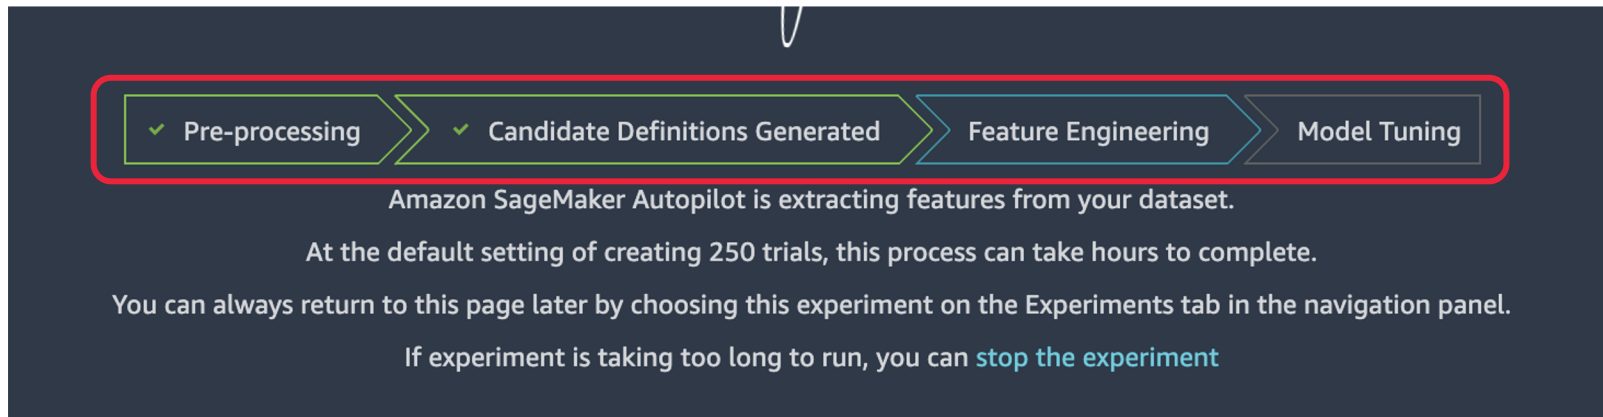
Job completion criteria

Max. training job run time

Max. AutoML job runtime

Specify input data

Monitor Progress in Amazon SageMaker Studio



The screenshot shows a progress bar with four steps: Pre-processing, Candidate Definitions Generated, Feature Engineering, and Model Tuning. The first two steps are marked with green checkmarks and are enclosed in a red rounded rectangle. The text below the progress bar states: "Amazon SageMaker Autopilot is extracting features from your dataset. At the default setting of creating 250 trials, this process can take hours to complete. You can always return to this page later by choosing this experiment on the Experiments tab in the navigation panel. If experiment is taking too long to run, you can [stop the experiment](#)".

✓ Pre-processing

✓ Candidate Definitions Generated

Feature Engineering

Model Tuning

Amazon SageMaker Autopilot is extracting features from your dataset.

At the default setting of creating 250 trials, this process can take hours to complete.

You can always return to this page later by choosing this experiment on the Experiments tab in the navigation panel.

If experiment is taking too long to run, you can [stop the experiment](#)

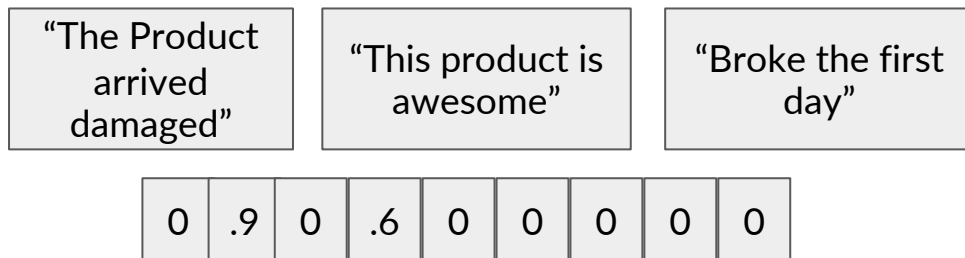
API → *DescribeAutoMLJob*

Generated Code for Feature Engineering

- SageMaker Autopilot automatically performs data exploration and prepares the data for the problem type

Example(s):


Transform text features
using
MultiColumnTfidfVectorizer



- SageMaker Autopilot will automatically tune `MultiColumnTfidfVectorizer` parameters


Bag-of-Words: Text as Vectors

I love this movie! It's sweet, but with satirical humor. The dialogue is great and the adventure scenes are fun... It manages to be whimsical and romantic while laughing at the conventions of the fairy tale genre. I would recommend it to just about anyone. I've seen it several times, and I'm always happy to see it again whenever I have a friend who hasn't seen it yet!



Term	Term Count
it	6
I	5
the	4
to	3
and	3
seen	2
yet	1
...	...

Text Mining: Measuring Word Importance



Statistical measure of word importance in corpus

↑ proportionally to the number of times a word appears in document

↓ by the frequency of the word in the corpus

tf-idf: term frequency-inverse document frequency

Computing Term Frequency (TF)

t	term
d	document
D	corpus

$$\text{tf}(t, d) = \frac{f_{t,d}}{\sum_{t' \in d} f_{t',d}}$$

Computing Inverse Document Frequency (IDF)


t	term
d	document
D	corpus

$$\text{idf}(t, D) = \log \left(\frac{|D|}{|\{d \in D : t \in d\}|} \right)$$

Putting It All Together: TF-IDF

$$\text{tf-idf}(t, d, D) = \text{tf}(t, d) \cdot \text{idf}(t, D)$$

I love this movie! It's sweet, but with satirical humor. The dialogue is great and the adventure scenes are fun... It manages to be whimsical and romantic while laughing at the conventions of the fairy tale genre. I would recommend it to just about anyone. I've seen it several times, and I'm always happy to see it again whenever I have a friend who hasn't seen it yet!



Term	TF / IDF
it	0.06
I	0.05
the	0.01
to	0.03
and	0.03
seen	0.04
yet	0.01
...	...

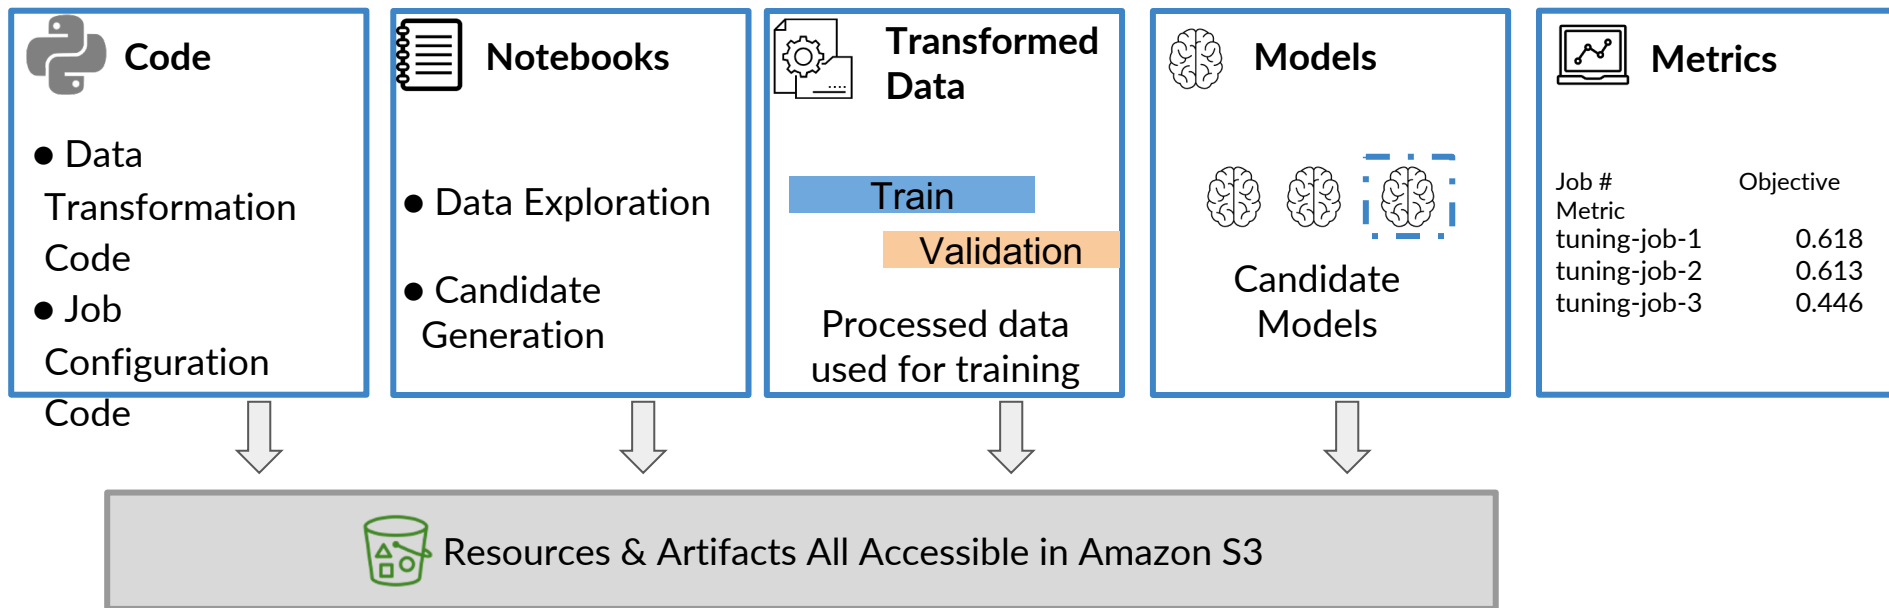
AutoML

with
Amazon SageMaker Autopilot

Evaluating Output




SageMaker Autopilot Generates Resources & Artifacts



Model Candidate Pipelines


A model candidate pipeline is composed of

- the feature engineering code (i.e. dpp0.py)
- and an algorithm (i.e. XGBoost).



Executing Model Candidate Pipelines

Model Candidate Pipelines



```
automl_interactive_runner.fit_data_transformers(parallel_jobs=7)
```

Screen recording (Shelbee to insert)

Model Hosting

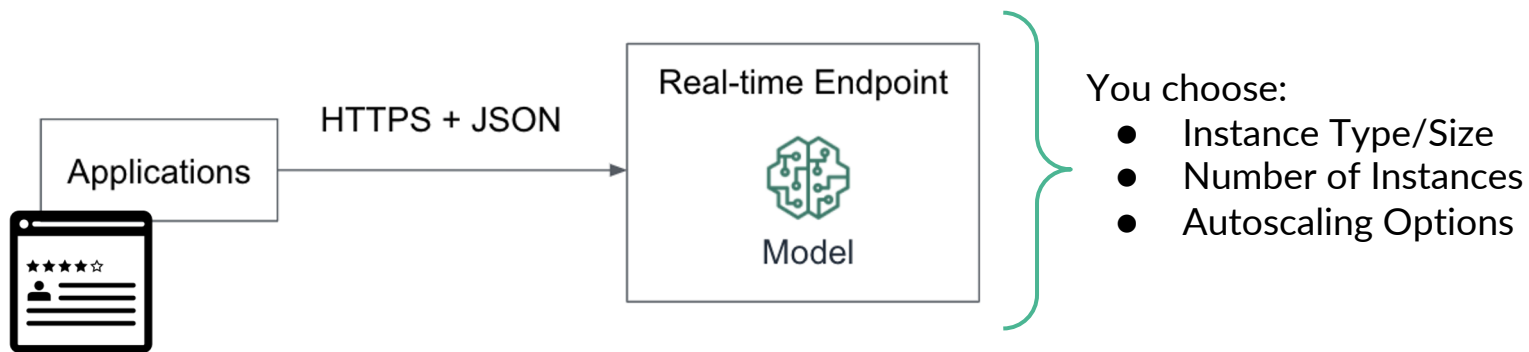
Introduction



Host a Model Endpoint

Deploy the model to serve predictions in real-time.

- Optimized for **low latency** of model predictions
- Example: As product reviews are coming in through various online channels, you want to predict the sentiment



Deploy Inference Pipeline

```
pipeline_model.deploy(initial_instance_count=1,  
                      instance_type='ml.m5.2xlarge',  
                      endpoint_name=pipeline_model.name,  
                      wait=True)
```

Congratulations! Now you could visit the sagemaker [endpoint console page](#) to find the deployed endpoint (it'll take a few minutes to be in service).

The `PipelineModel` has multiple containers of the following:

- **Data Transformation Container:** a container built from the model we selected and trained during the data transformer sections
- **Algorithm Container:** a container built from the trained model we selected above from the best HPO training job.
- **Inverse Label Transformer Container:** a container that converts numerical intermediate prediction value back to non-numerical label value.

Inference Pipeline

