

# Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see <https://creativecommons.org/licenses/by-sa/2.0/legalcode>



DeepLearning.AI



# Welcome to Course 3

---



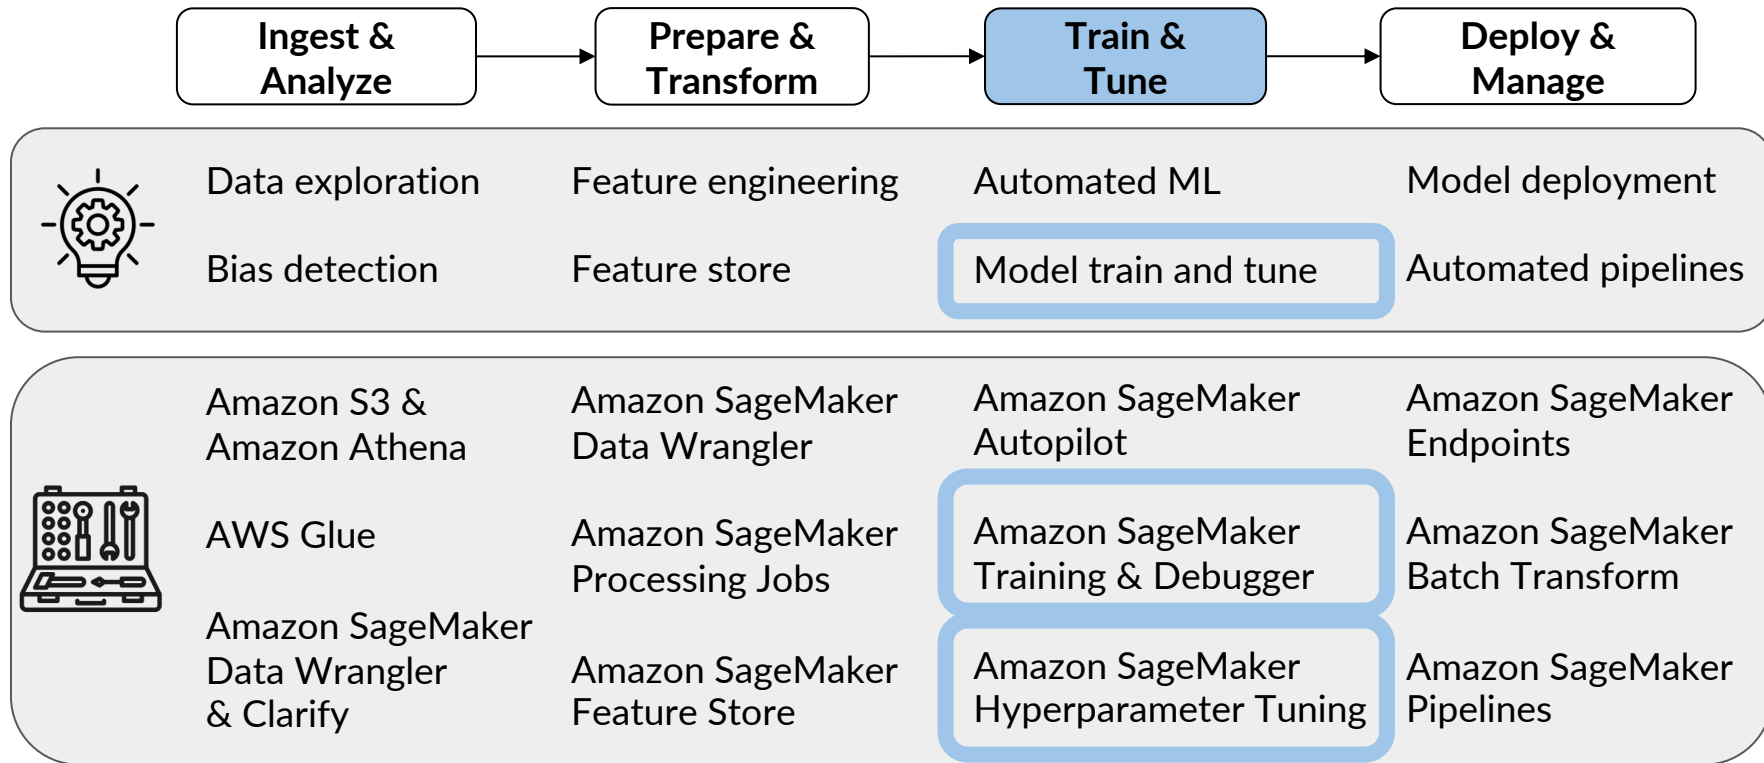
DeepLearning.AI



# Advanced Model Training

---

# Machine Learning Workflow



# Model Tuning

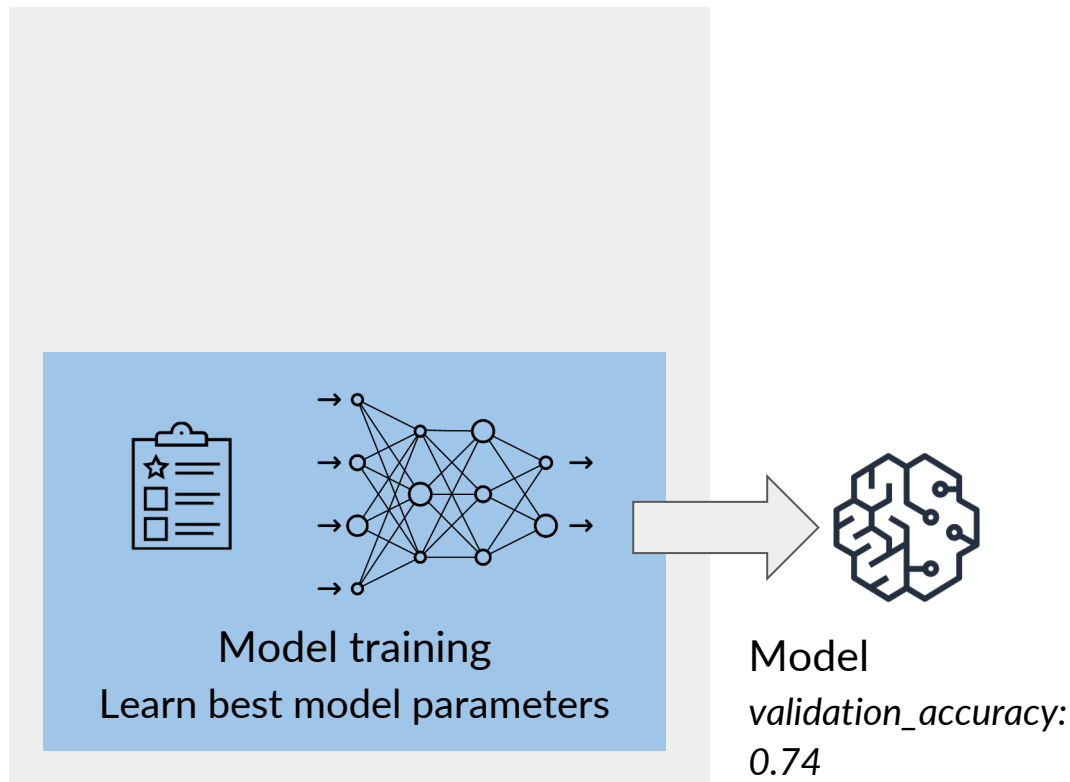


# Model Tuning

## Model parameters



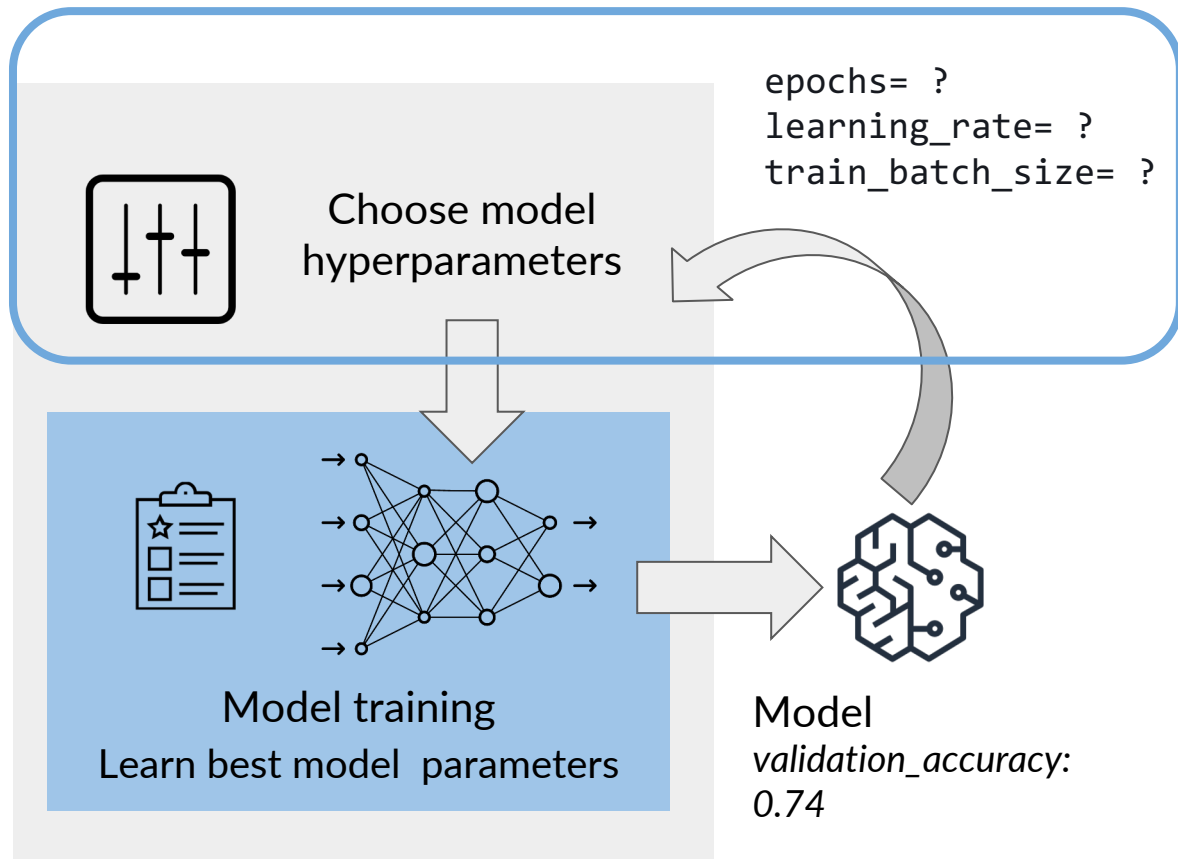
*roberta-base*  
125M parameters



# Model Tuning

Model hyperparameters

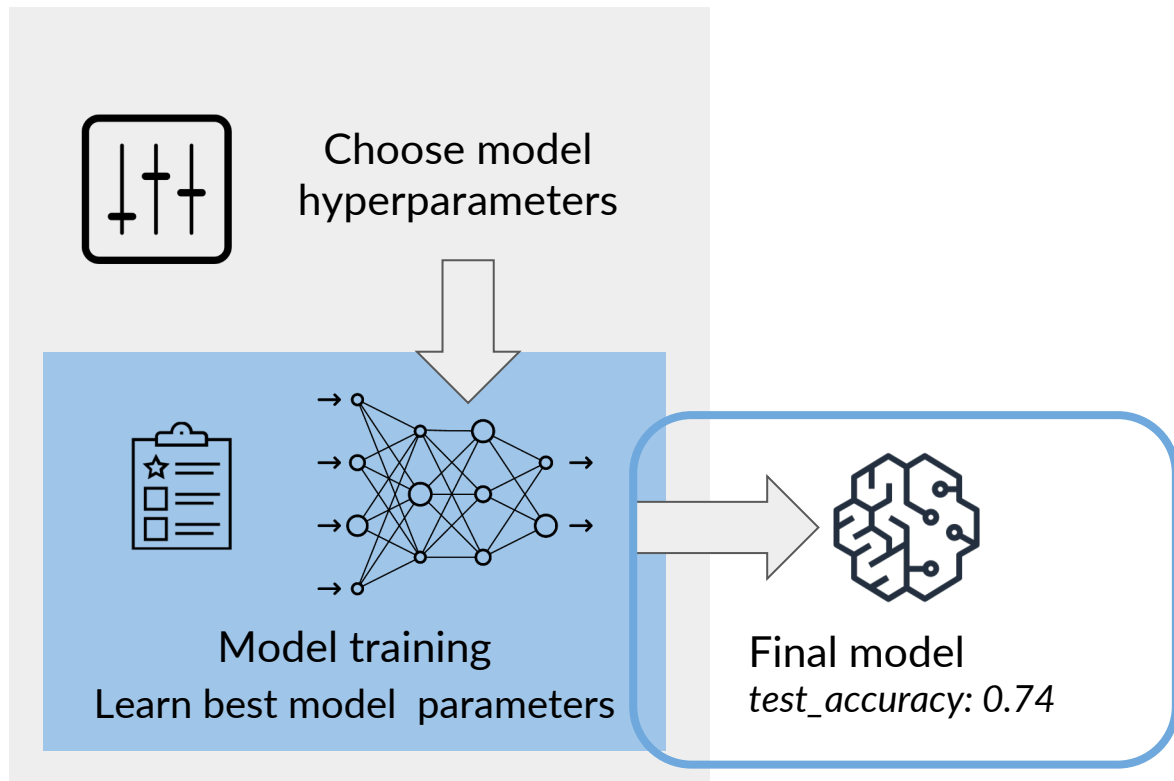
Model parameters



# Model Evaluation

"If you can't measure it,  
you can't improve it."

-- Peter Drucker






# Manual vs. Automatic Model Tuning



Manual tuning



Automatic  
model tuning

# Popular Algorithms for Automatic Model Tuning





Automatic  
model tuning

- Grid search
- Random search
- Bayesian optimization
- Hyperband



# Grid Search

- Define sets of hyperparameters
- Test **every** combination
- Select the best performing hyperparameters



# Grid Search



- Define sets of hyperparameters
- Test **every** combination
- Select the best performing hyperparameters



- + Explores all combinations
- + Works for small number of parameters
- Time-consuming
- Doesn't scale to large numbers of parameters



# Random Search

- Define sets of hyperparameters
  - Define search space & stop criteria
  - Test **random** combinations within search space
  - Select the best performing hyperparameters
- 



# Random Search



- Define sets of hyperparameters
  - Define search space & stop criteria
  - Test **random** combinations within search space
  - Select the best performing hyperparameters
- 



- + Faster compared to grid search
- Might miss better performing hyperparameters



# Bayesian Optimization

- Treat HPT like a **regression** problem (surrogate model)
- Start from random hyperparameters
- Narrow down search space around better performing hyperparameters




# Bayesian Optimization

- Treat HPT like a **regression** problem (surrogate model)
- Start from random hyperparameters
- Narrow down search space around better performing hyperparameters





- + More efficient in finding best hyperparameters
- Requires sequential execution
- Might get stuck in local minima






# Hyperband


- **Bandit-based** approach
- Start from random hyperparameters
- Explore sets of hyperparameters for few iterations
- Choose best and explore longer
- Repeat until `max_iterations` reached or one candidate left



# Hyperband



- **Bandit-based** approach
- Start from random hyperparameters
- Explore sets of hyperparameters for few iterations
- Choose best and explore longer
- Repeat until `max_iterations` reached or one candidate left

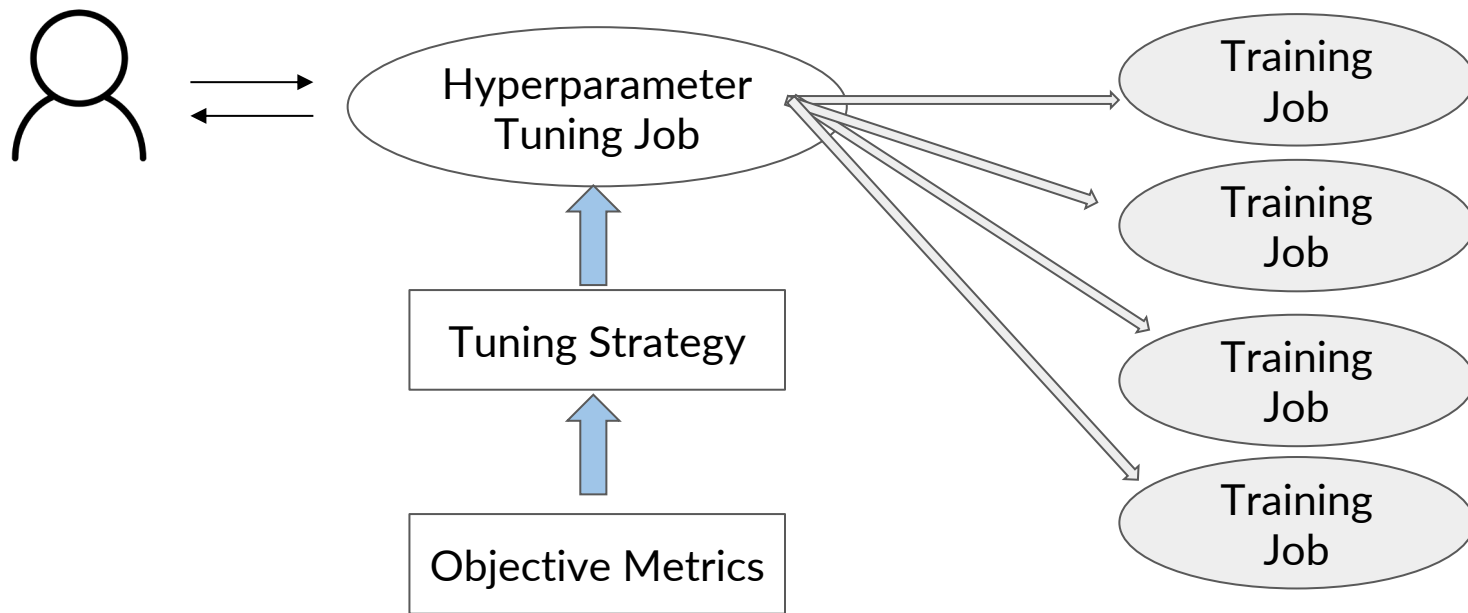


- + Spends time efficiently (explore-exploit theory)
- Might discard good candidates early that converge slowly

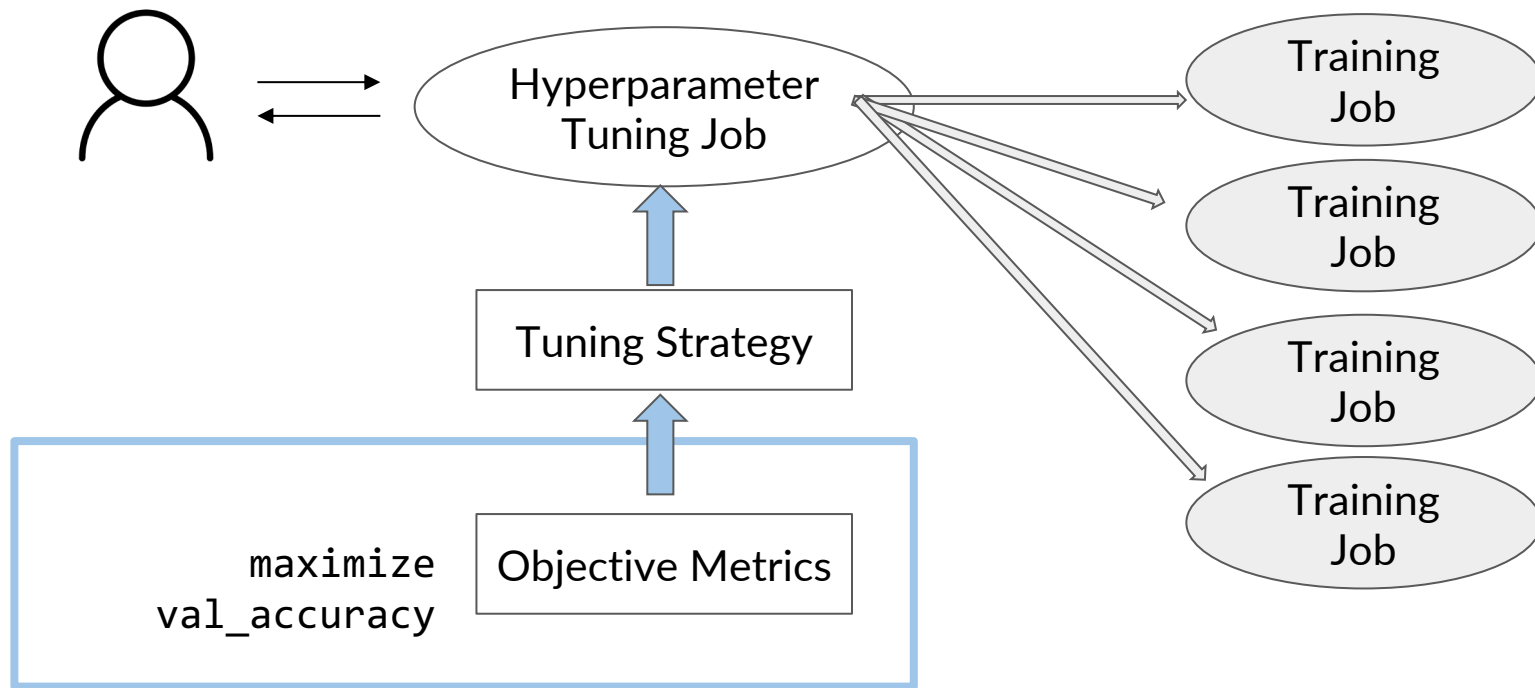
# Tune a BERT- based Text Classifier



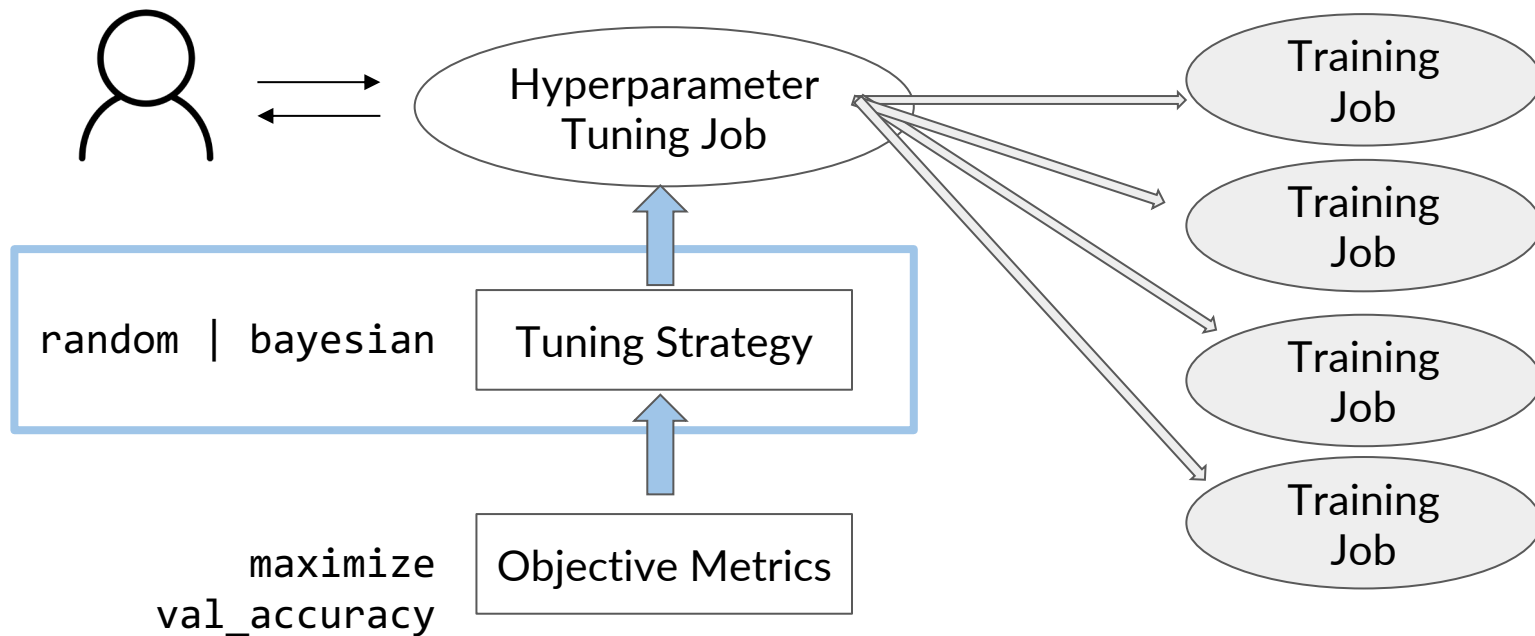
# Amazon SageMaker Hyperparameter Tuning (HPT)



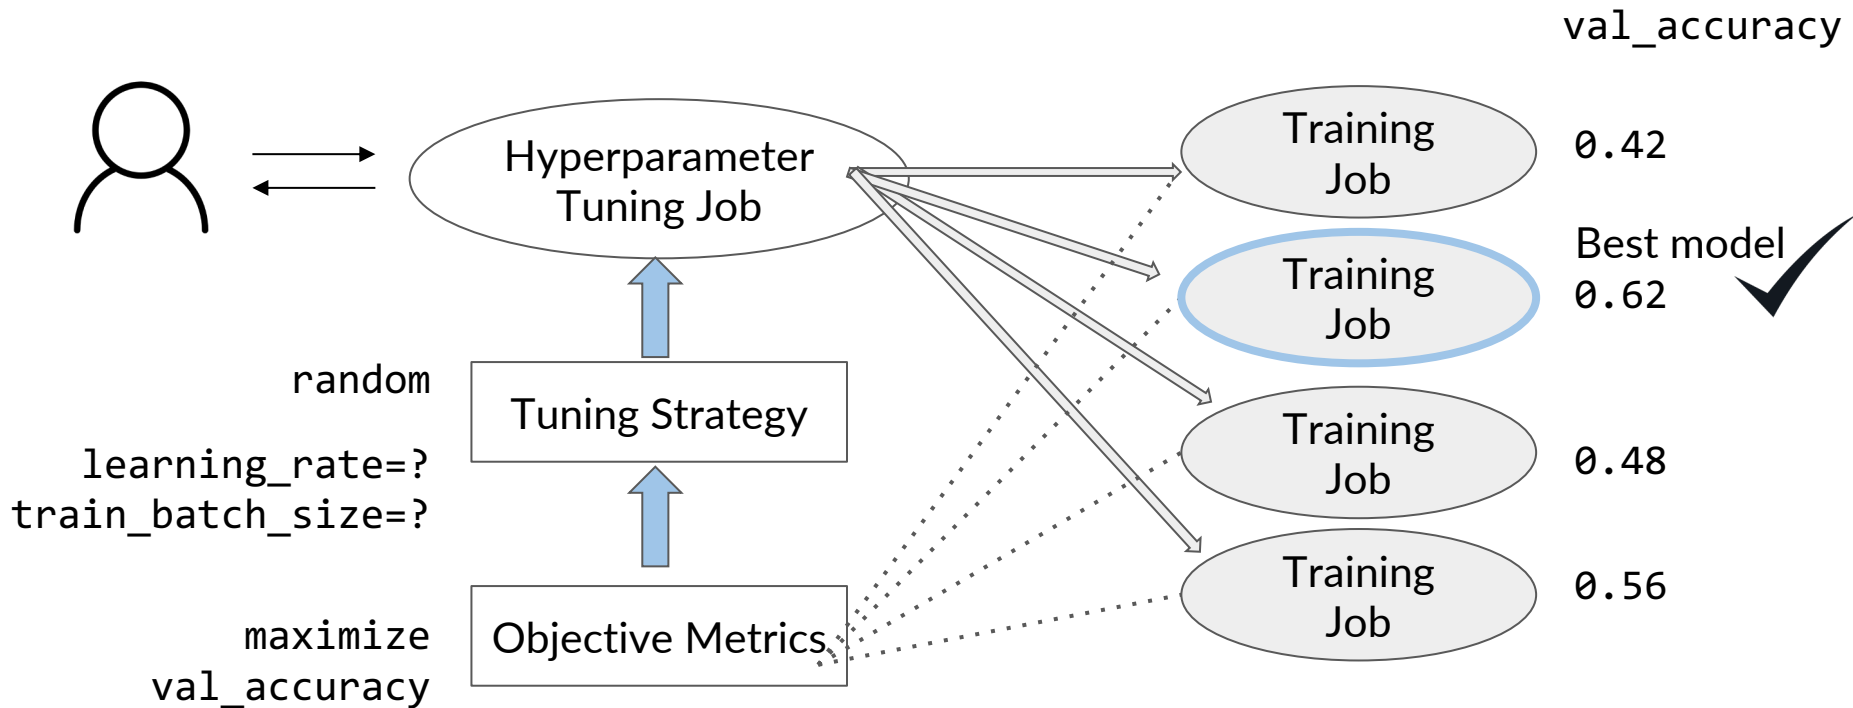
# Amazon SageMaker Hyperparameter Tuning (HPT)



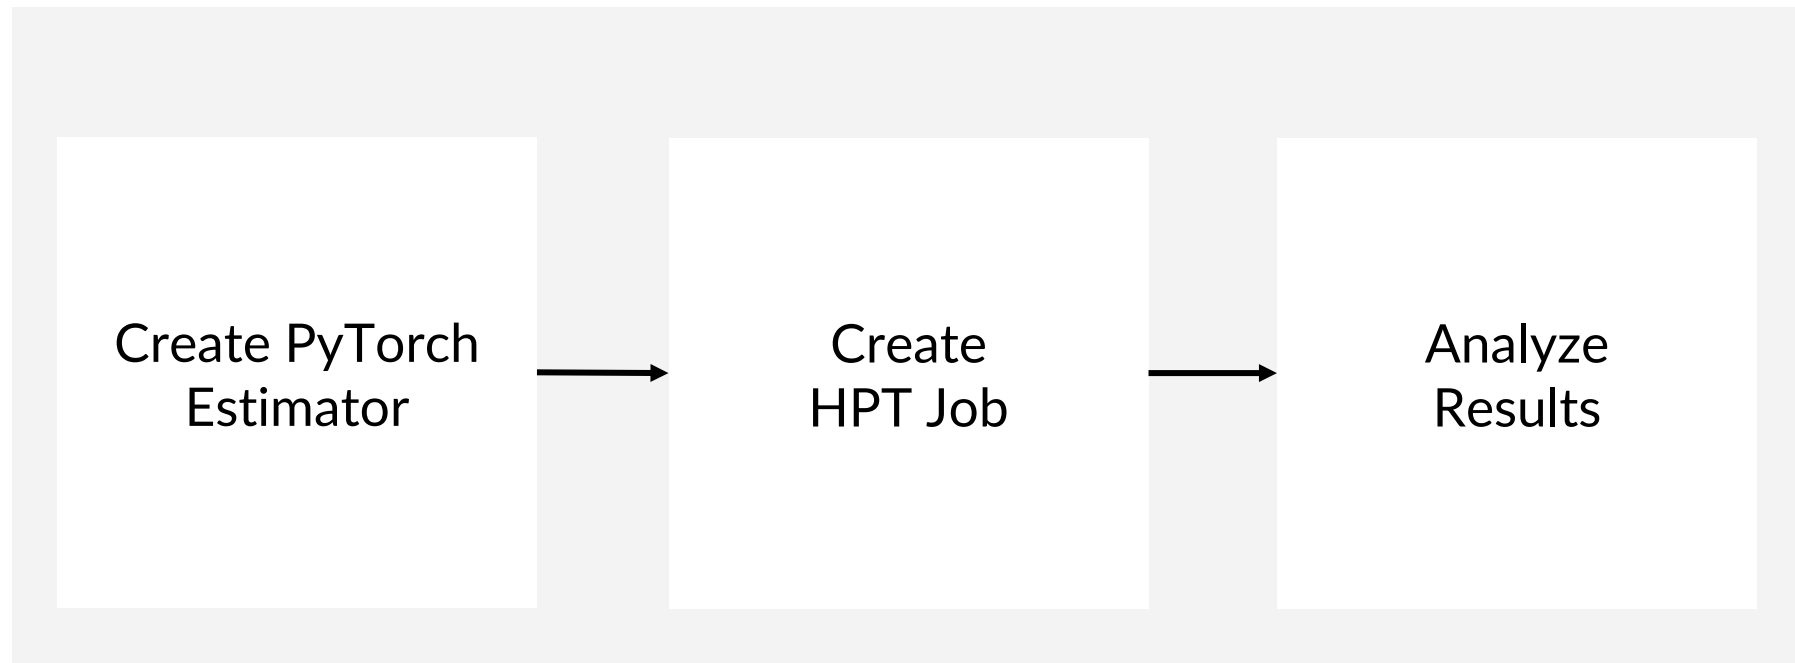
# Amazon SageMaker Hyperparameter Tuning (HPT)



# Tune BERT Text Classifier

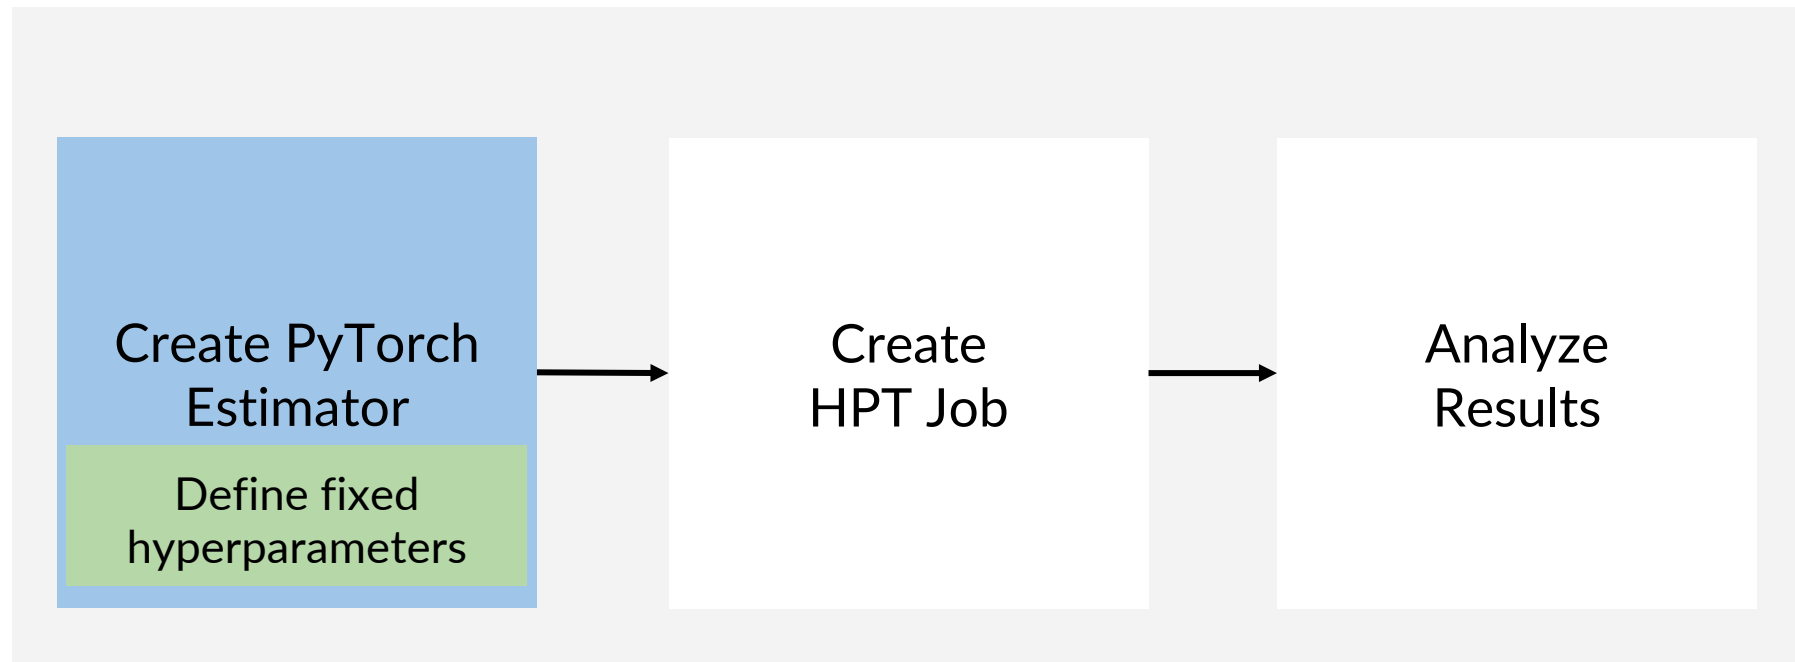


# Steps





# Steps



# Define Fixed Hyperparameters

```
hyperparameters={  
    'epochs': 3,  
    'train_steps_per_epoch': 50,  
    'validation_batch_size': 64,  
    'validation_steps_per_epoch': 50,  
    'freeze_bert_layer': False,  
    'seed': 42,  
    'max_seq_length': 64,  
    'backend': 'gloo',  
    'run_validation': True,  
    'run_sample_predictions': False  
}
```

Create PyTorch estimator

Define fixed  
hyperparameters

# Create PyTorch Estimator

Create PyTorch estimator

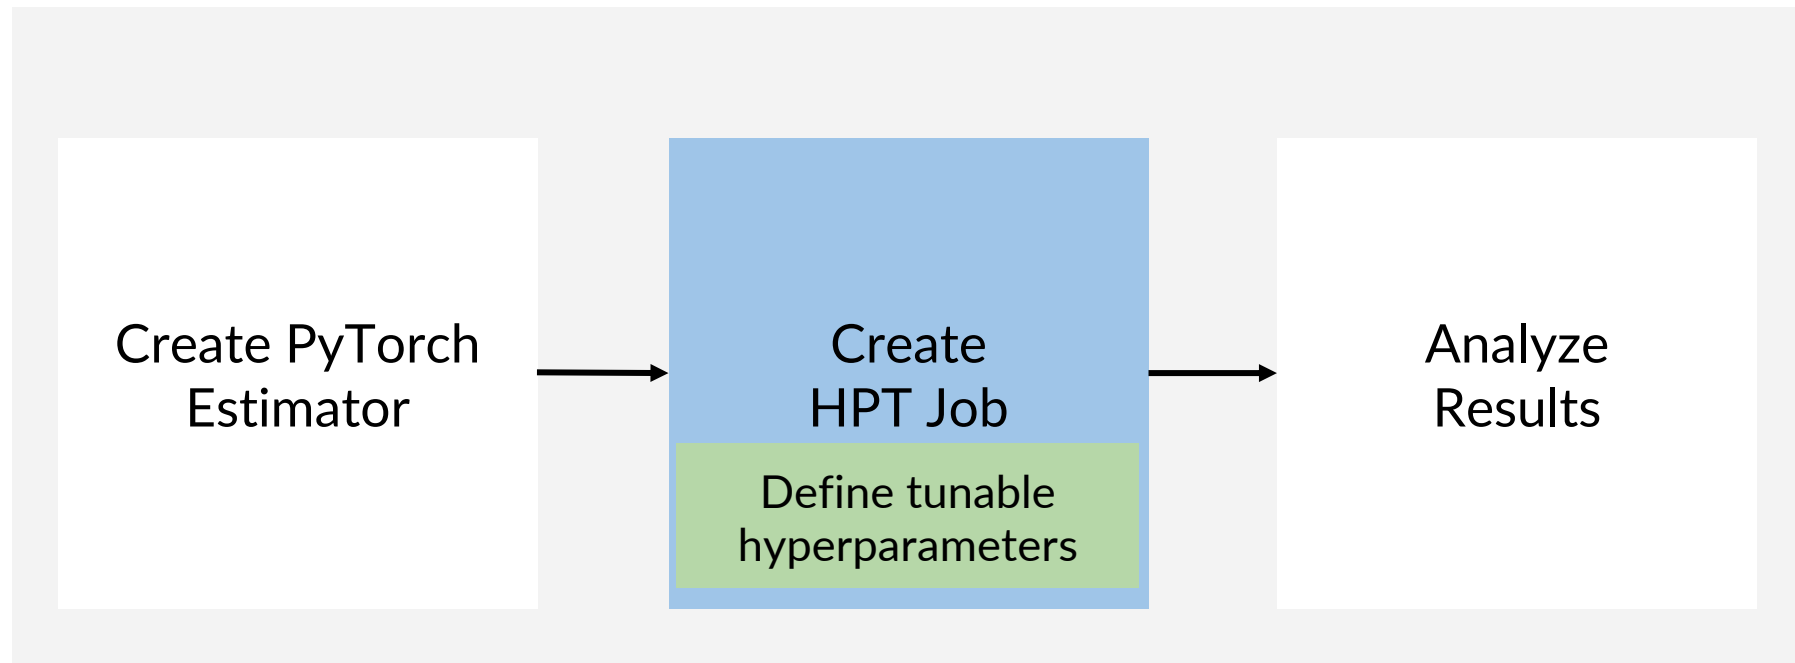
Define fixed  
hyperparameters

```
from sagemaker.pytorch import PyTorch as PyTorchEstimator
```

```
estimator = PyTorchEstimator(  
    entry_point='train.py',  
    ...  
    hyperparameters=hyperparameters,  
  
)
```

Fixed hyperparameters  
in estimator

# Steps



# Define Tunable Hyperparameters

Create HPT job

Define tunable  
hyperparameters

```
from sagemaker.tuner import CategoricalParameter
from sagemaker.tuner import ContinuousParameter
from sagemaker.tuner import IntegerParameter
```

Specify parameters

```
hyperparameter_ranges = {
    'learning_rate': ContinuousParameter(0.00001, 0.00005,
scaling_type='Linear'),
    'train_batch_size': CategoricalParameter([128, 256]),
}
```

Specify hyperparameter types

Specify ranges

# How to Choose Hyperparameter Types

## Categorical

'train\_batch\_size':

CategoricalParameter([128, 256])

'freeze\_bert\_layer':

CategoricalParameter([True, False])

# How to Choose Hyperparameter Types

## Categorical

'train\_batch\_size':

CategoricalParameter([128, 256])

'freeze\_bert\_layer':

CategoricalParameter([True, False])

## Integer

'train\_batch\_size':

IntegerParameter(16, 1024, scaling\_type='Logarithmic')

If you need to  
explore large  
ranges quickly

# How to Choose Hyperparameter Types

## Categorical

'train\_batch\_size':

CategoricalParameter([128, 256])

'freeze\_bert\_layer':

CategoricalParameter([True, False])

## Integer

'train\_batch\_size':

IntegerParameter(16, 1024, scaling\_type='Logarithmic')

## Continuous

'learning\_rate':

ContinuousParameter(0.00001, 0.00005, scaling\_type='Linear')



# Create Amazon SageMaker HPT job

Create HPT job

Define tunable  
hyperparameters

```
from sagemaker.tuner import HyperparameterTuner
```

```
tuner = HyperparameterTuner(  
    estimator=...,  
    hyperparameter_ranges=...,  
    objective_type=...,  
    objective_metric_name=...,  
    strategy=...,  
)
```

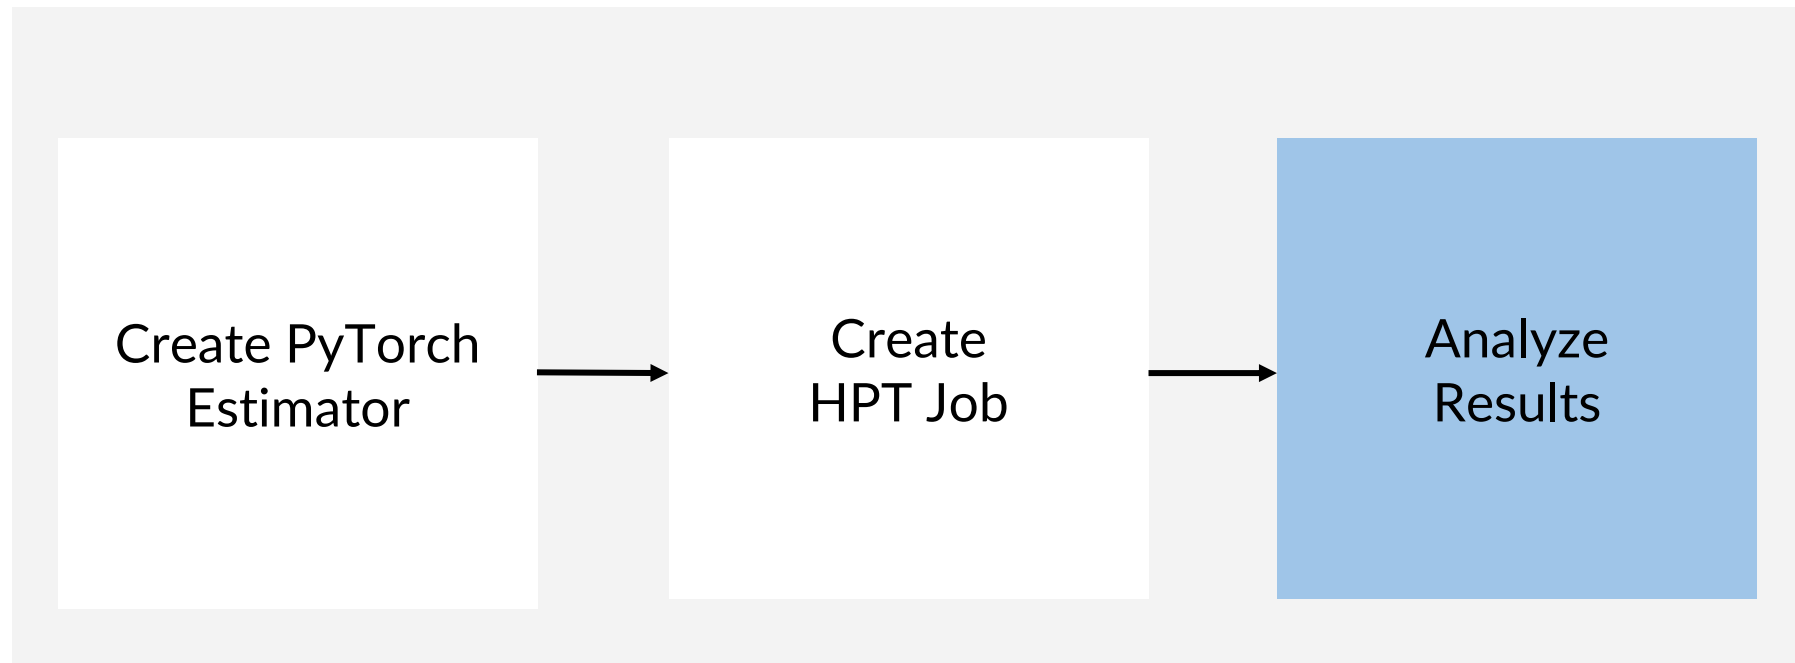
Pass in estimator

Configure  
Hyperparameter ranges

Run HPT job with .fit()

```
tuner.fit(inputs={...}, ...)
```

# Steps



# Analyze Results

Analyze Results

```
df_results = tuner.analytics().dataframe()
```

# Analyze Results

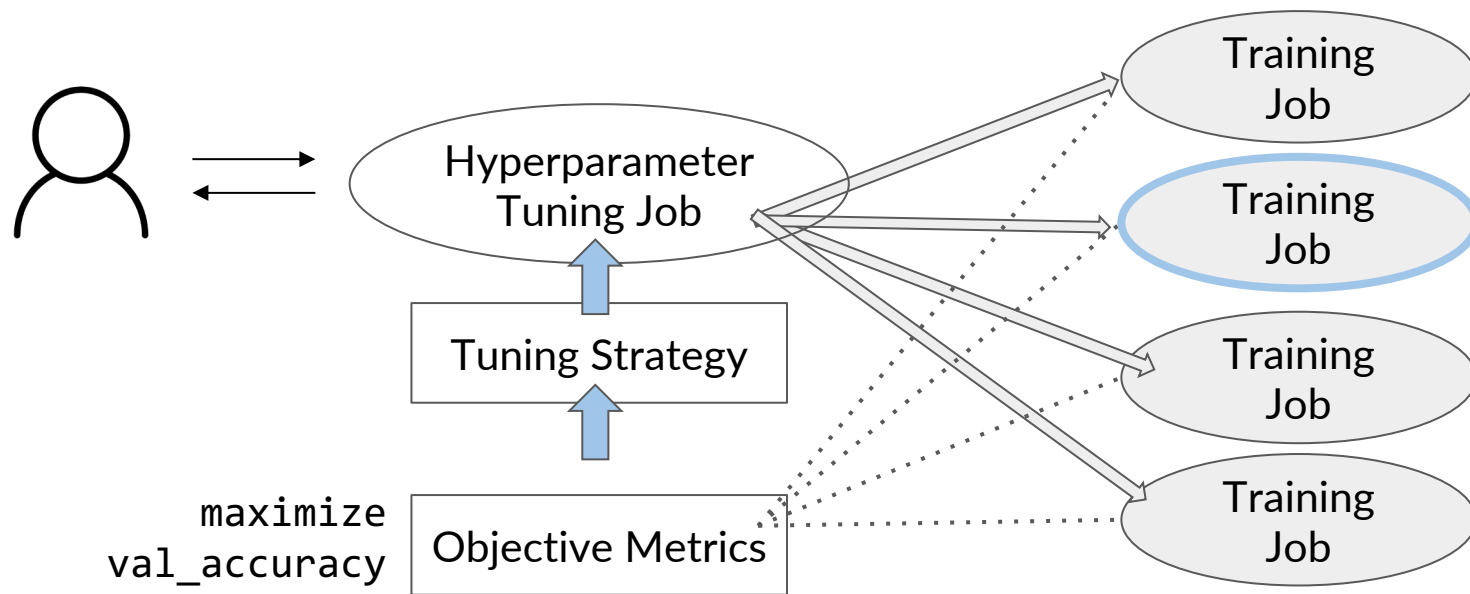
Analyze Results

```
df_results = tuner.analytics().dataframe()
```

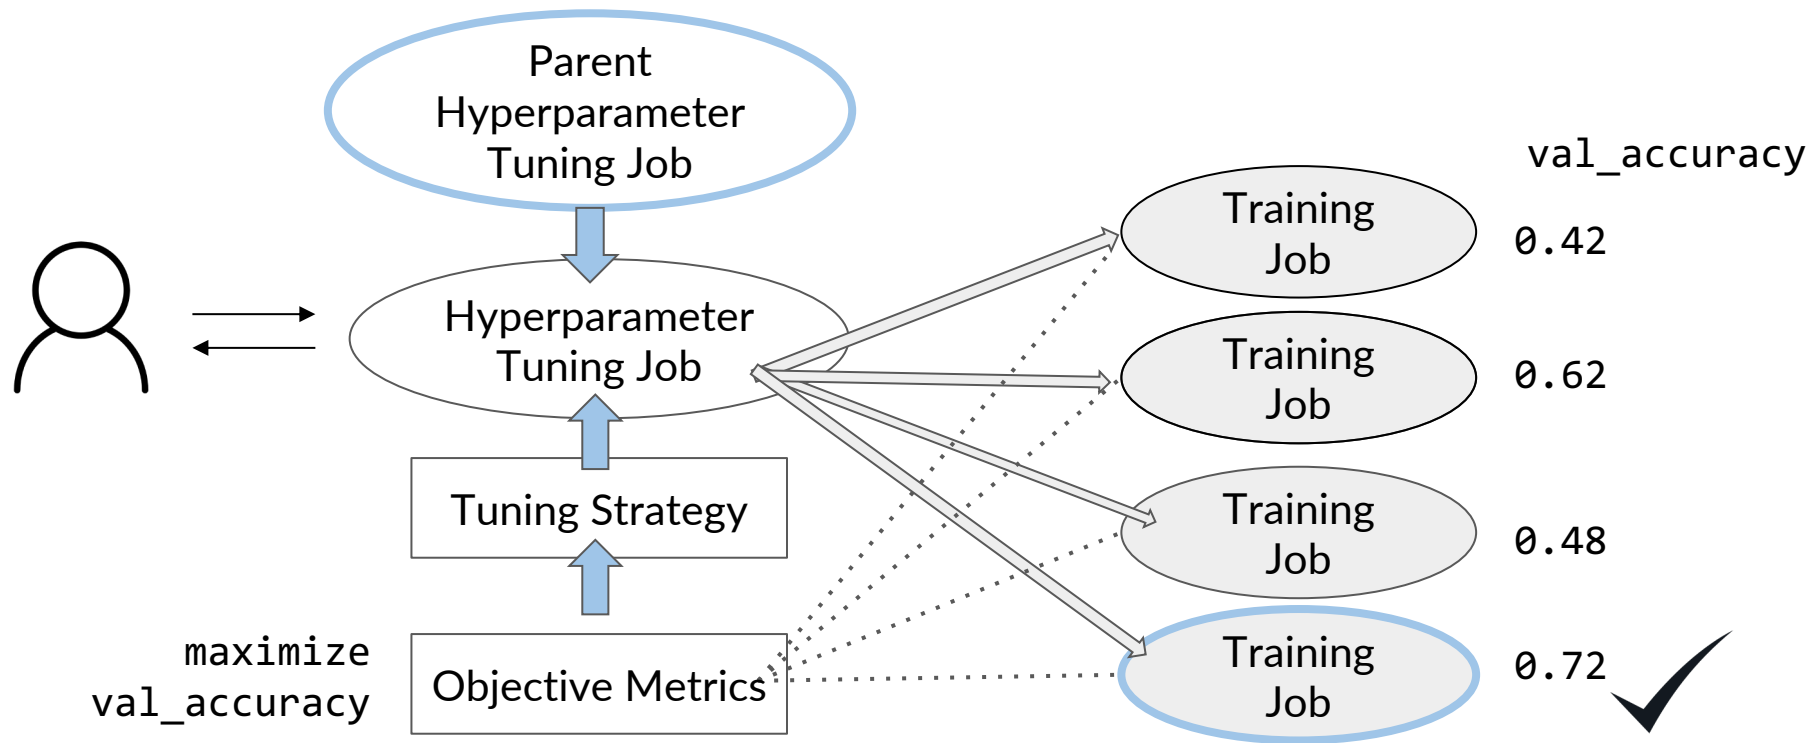
learning_rate	train_batch_size	TrainingJobName	TrainingJobStatus	FinalObjectiveValue
0.000021	"128"	pytorch-training-210225-1535-001-71394bc3	Completed	44.939999
0.000035	"128"	pytorch-training-210225-1535-002-cf437bad	Completed	41.580002



# Warm Start HPT Job



# Warm Start HPT Job



# Warm Start HPT Job

- IDENTICAL\_DATA\_AND\_ALGORITHM
  - Same input data and training data
  - Update hyperparameter tuning ranges and maximum number of training jobs
- TRANSFER\_LEARNING
  - Updated training data and different version of training algorithm

# Configure Warm Start

```
from sagemaker.tuner import WarmStartConfig
from sagemaker.tuner import WarmStartTypes
```

```
warm_start_config = WarmStartConfig(
    warm_start_type=WarmStartTypes.IDENTICAL_DATA_AND_ALGORITHM,
    parents=<PARENT_TUNING_JOB_NAME>)
```

IDENTICAL\_DATA\_AND\_ALGORITHM  
or TRANSFER\_LEARNING

```
tuner = HyperparameterTuner(
    ...
    warm_start_config=warm_start_config)
```

Specify parent tuning job


```
tuner.fit(...)
```

Pass warm start config  
in HyperparameterTuner



# Best Practices - SageMaker HyperParameter Tuning

- Select a small number of hyperparameters
- Select a small range for hyperparameters
- Enable warm start
- Enable early stop to save tuning time and costs
- Select a small number of concurrent training jobs



# Best Practices - Monitoring Training Resources

- Right size compute resources
- Requires empirical testing
- Amazon CloudWatch Metrics
- Insights from Amazon SageMaker Debugger

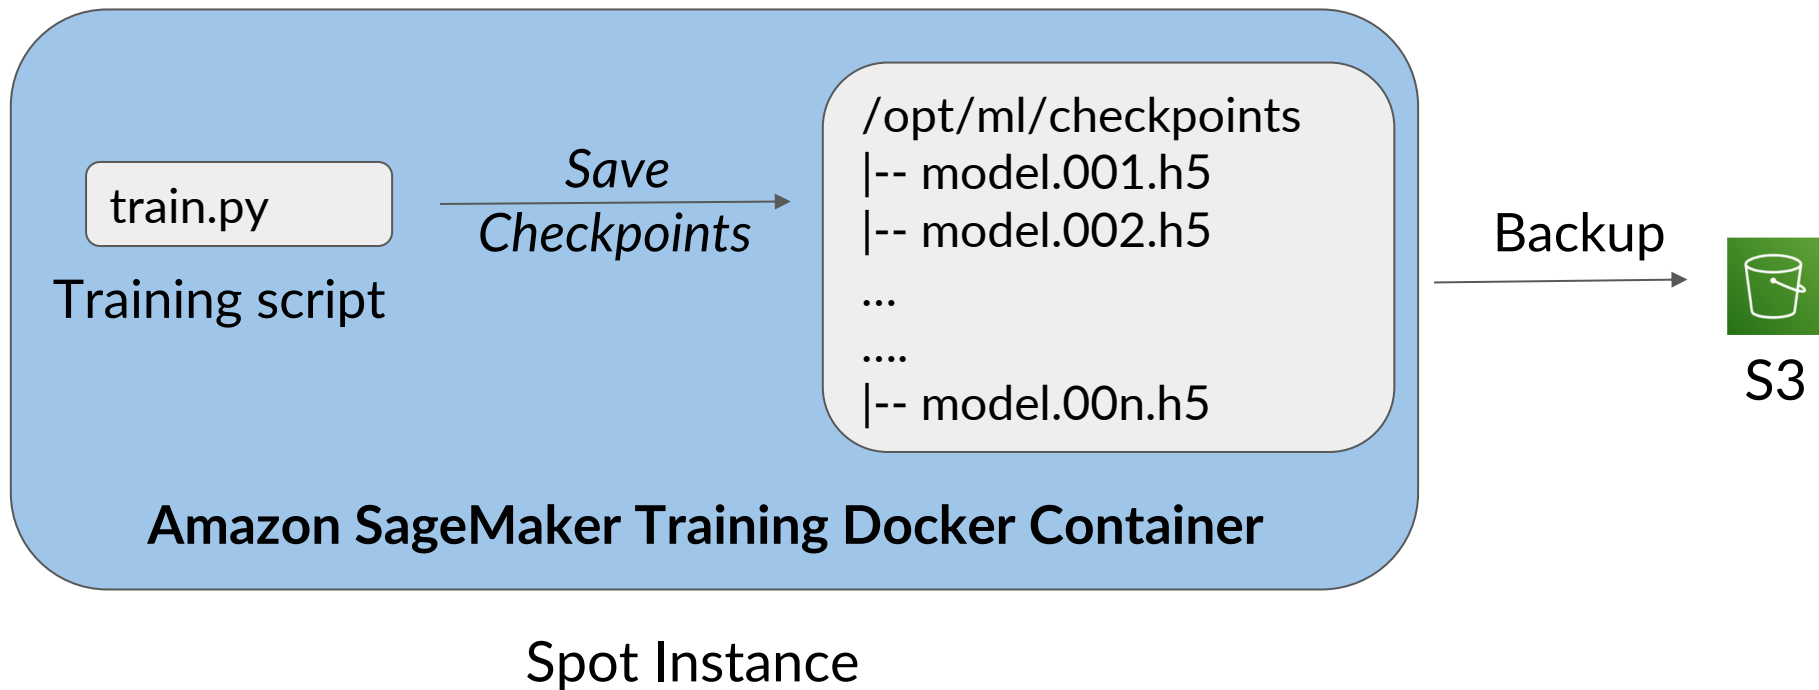
# Checkpointing



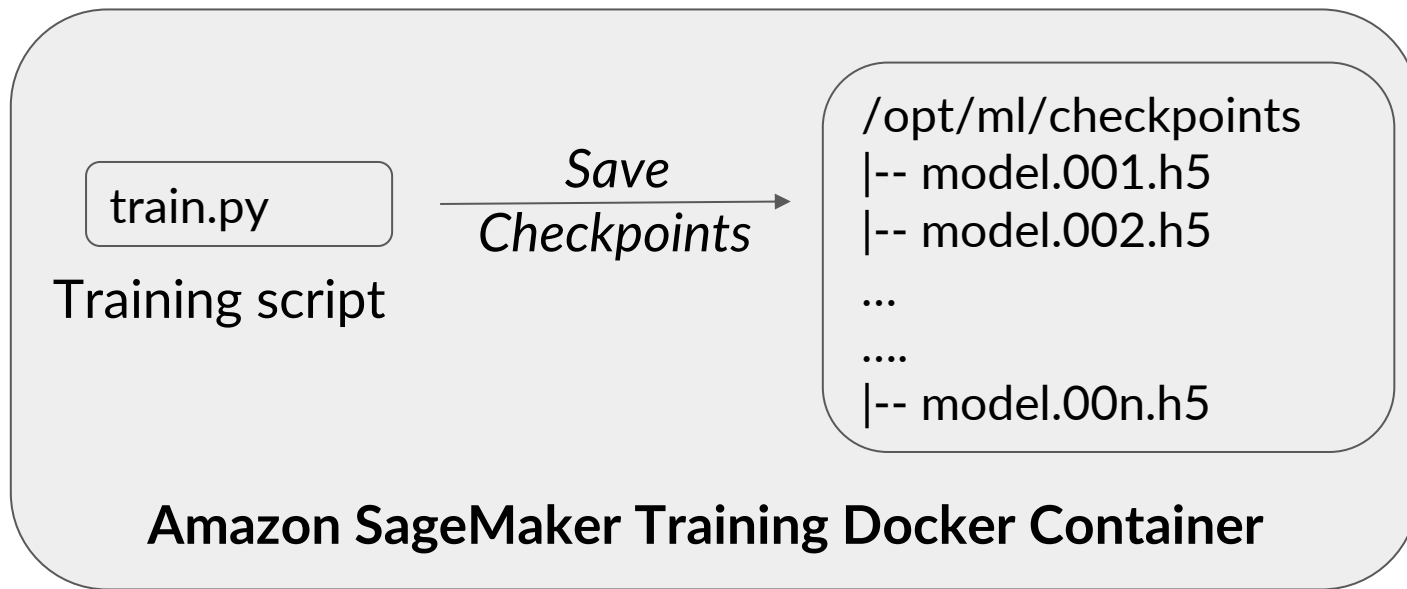
# Machine Learning Checkpointing

- Save state of ML models during training
- Checkpoints : Snapshots of the model
  - Model architecture
  - Model weights
  - Training configurations
  - Optimizer
- Frequency and number of checkpoints

# Amazon SageMaker Managed Spot



# Amazon SageMaker Managed Spot



S3

Spot Instance (Terminated)

# Amazon SageMaker Spot Training

Spot Instance (Terminated)

train.py

Training script

Save  
Checkpoints

/opt/ml/checkpoints

-- model.001.h5

-- model.002.h5

...

....

-- model.00n.h5

Amazon SageMaker Training Docker Container

Spot Instance (New)



S3

Sync

# Distributed Training Strategies





# Challenges




Increased training data volume




Increased model size and complexity

# Distributed Training




Compute Node



Compute Cluster

# Distributed Training Strategies




Data Parallelism



Model Parallelism

# Distributed Training Strategies - Data Parallelism

- Training data split up
- Model replicated on all nodes



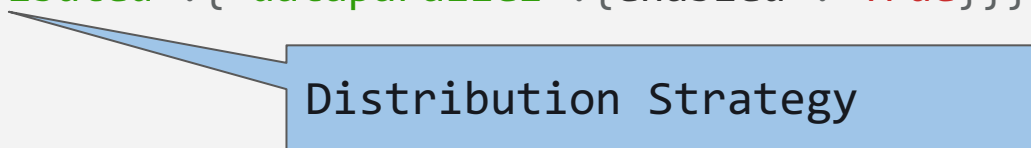
# Distributed Training Strategies - Model Parallelism

- Training data replicated
- Model split up on all nodes



# Amazon SageMaker Estimator

```
from sagemaker.pytorch import PyTorch
estimator = PyTorch(
    entry_point='train.py',
    role=sagemaker.get_execution_role(),
    framework_version='1.6.0',
    py_version='py3',
    instance_count=3,
    instance_type='ml.p3.16xlarge',
    distribution={'smdistributed':{'dataparallel':{'enabled': True}}}
)
estimator.fit()
```




Distribution Strategy

# Amazon SageMaker Estimator

```
from sagemaker.pytorch import PyTorch
estimator = PyTorch(
    entry_point='train.py',
    role=sagemaker.get_execution_role(),
    framework_version='1.6.0',
    py_version='py3',
    instance_count=3,
    instance_type='ml.p3.16xlarge',
    distribution={'smdistributed':{'modelparallel':{'enabled': True}}}
)
estimator.fit()
```

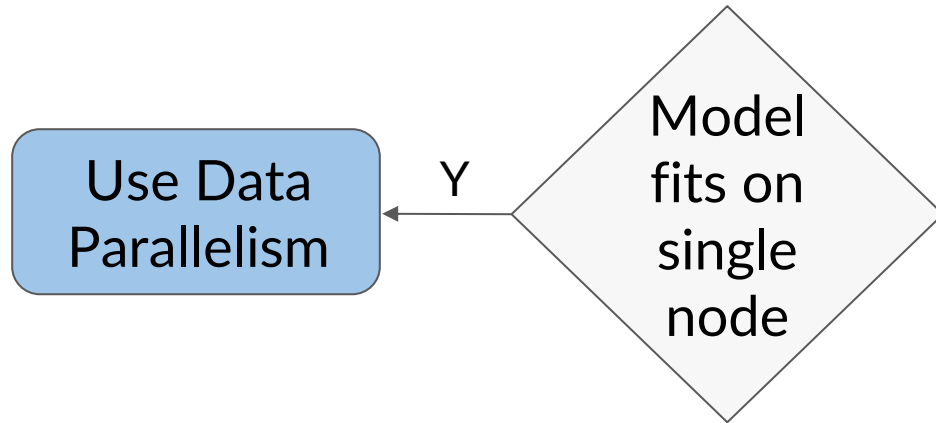
Model Parallel  
Distribution Strategy

# Choosing a Distribution Strategy







# Choosing a Distribution Strategy




# Choosing a Distribution Strategy




# Choosing a Distribution Strategy




# Choosing a Distribution Strategy




# Choosing a Distribution Strategy



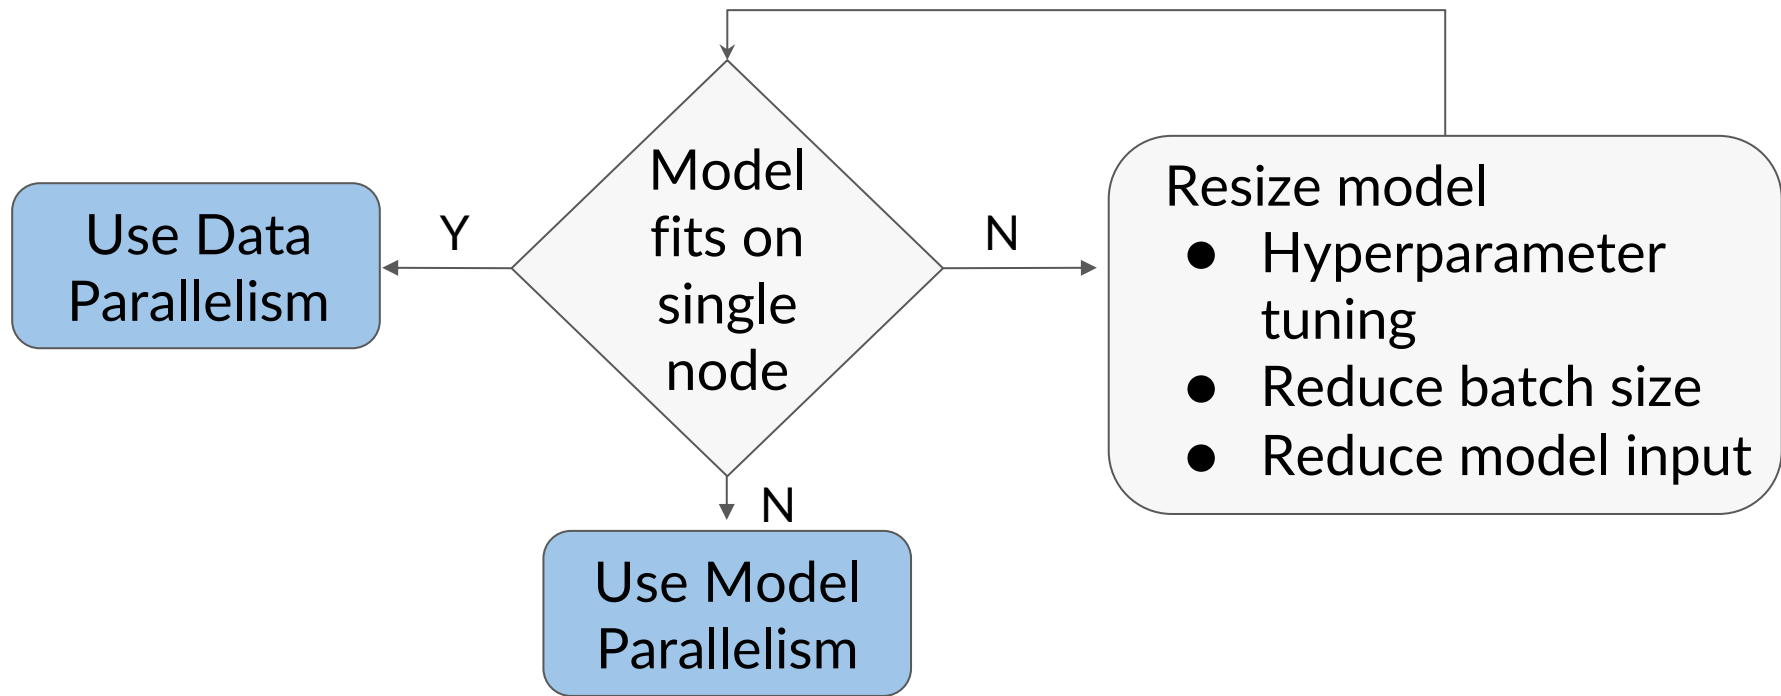
# Choosing a Distribution Strategy



# Choosing a Distribution Strategy



# Choosing a Distribution Strategy

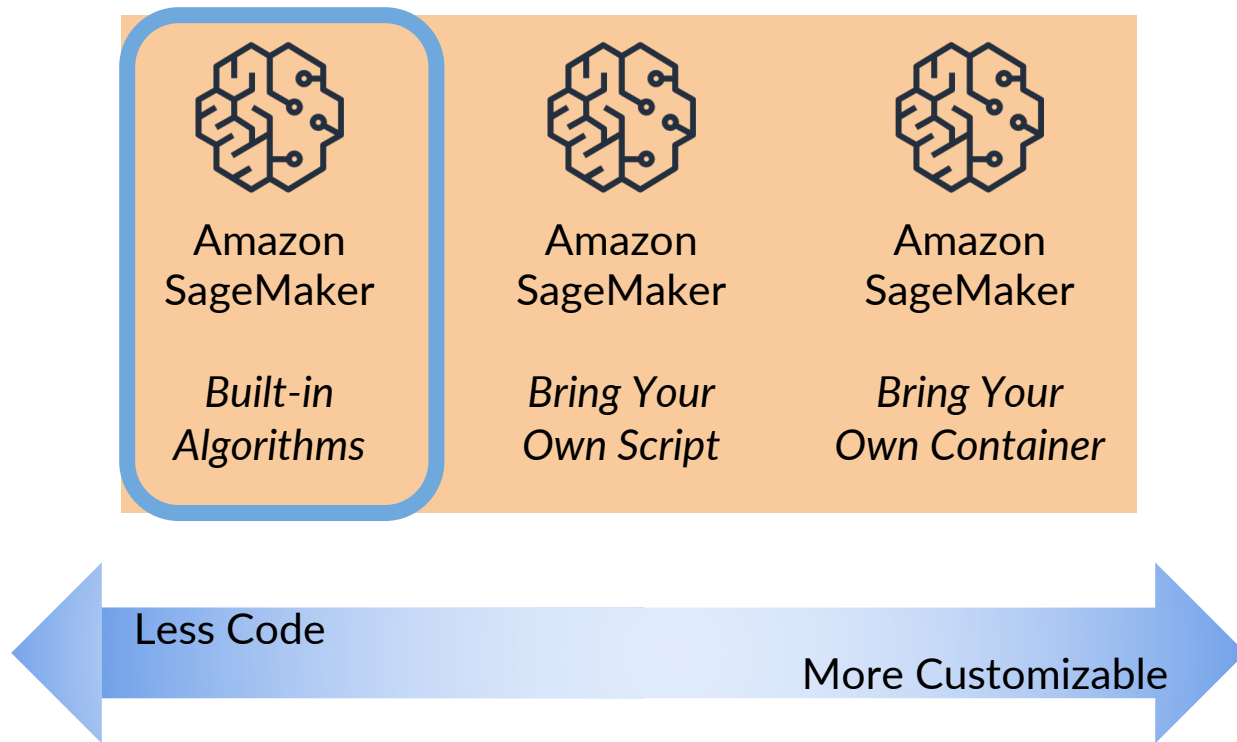




# Custom Algorithms with Amazon SageMaker



# Options on Amazon SageMaker

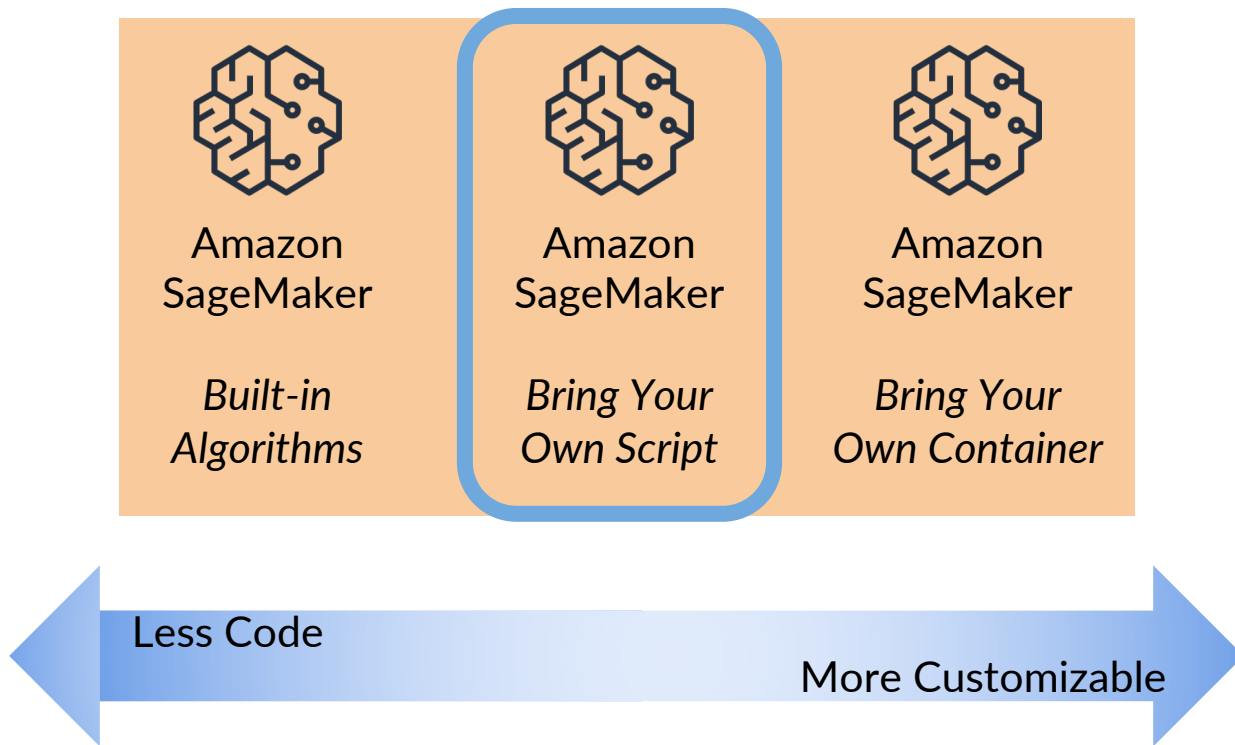


# Amazon SageMaker Estimator

```
estimator =  
sagemaker.estimator.Estimator(image_uri=image_uri, ...)  
estimator.set_hyperparameters(...)  
estimator.fit(...)
```

## Built-In Algorithms

# Options on Amazon SageMaker



# Amazon SageMaker Estimator

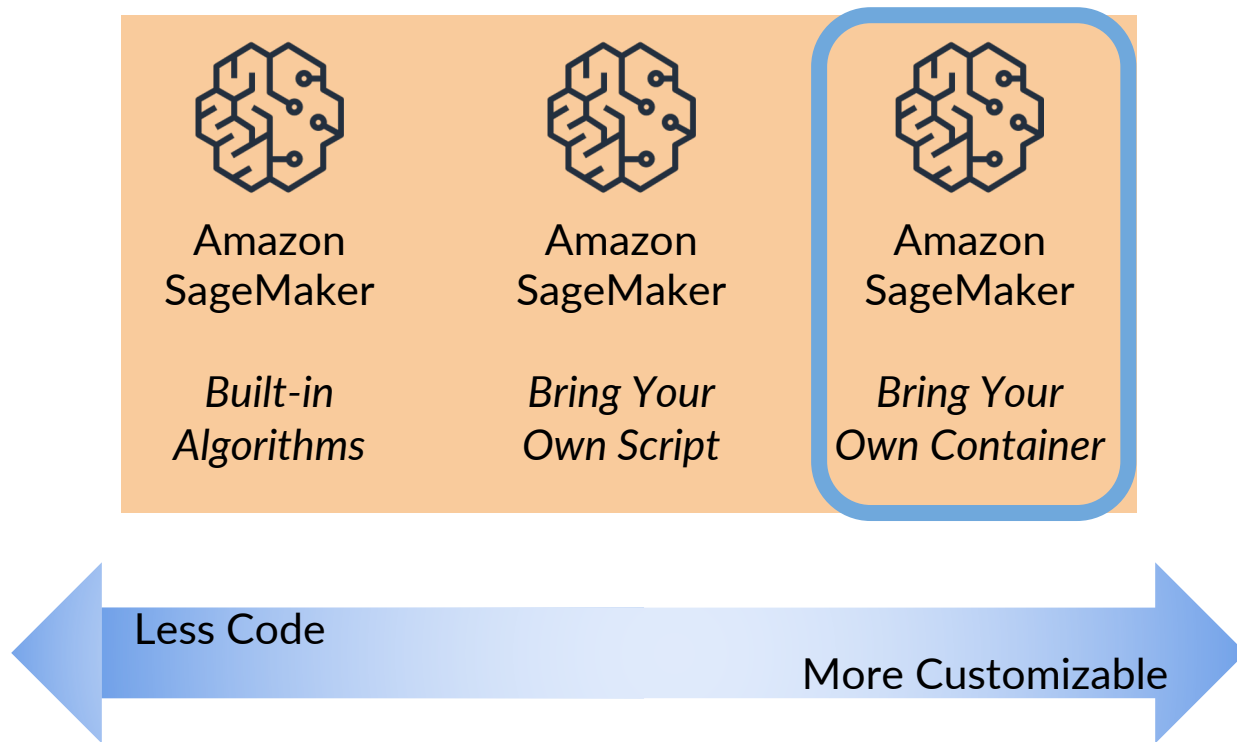
```
estimator =  
sagemaker.estimator.Estimator(image_uri=image_uri, ...)  
estimator.set_hyperparameters(...)  
estimator.fit(...)
```

**Built-In Algorithms**

```
from sagemaker.pytorch import PyTorch  
pytorch_estimator = PyTorch(  
    entry_point='train.py',  
    ...  
)
```

**Script Mode  
PyTorch Container**

# Training Options on Amazon SageMaker



# Bring Your Own Container - Steps

**Code**



**Containerize**




**Register with  
Amazon ECR**



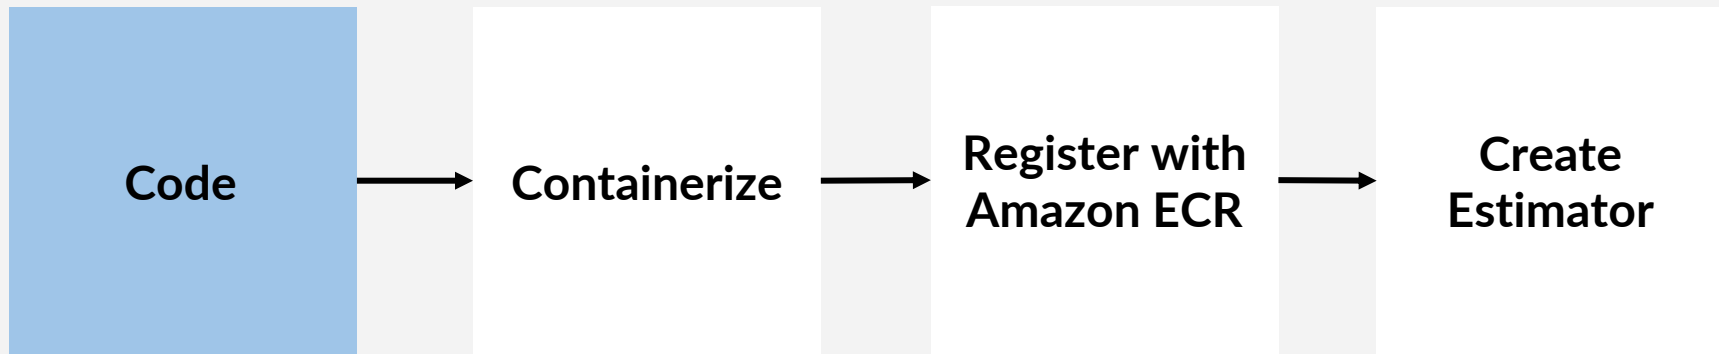
**Create  
Estimator**

# Bring Your Own Container - Steps





# Bring Your Own Container - Steps



- Algorithm
- Training
- Inference

# Bring Your Own Container - Steps

Code



Containerize

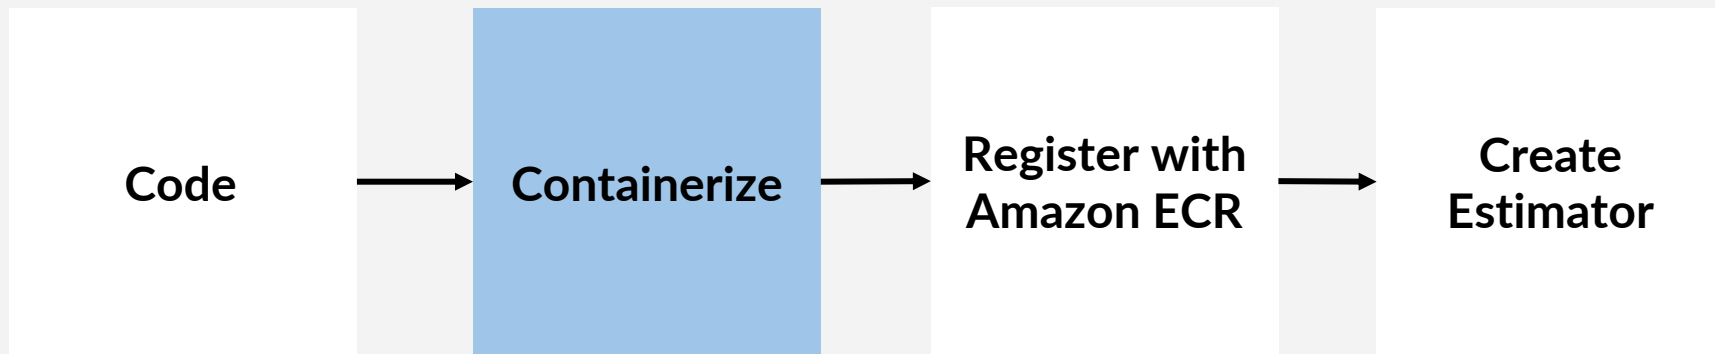


Register with  
Amazon ECR



Create  
Estimator

# Bring Your Own Container - Steps



- `algorithm_name=tf-custom-container-test`
- `docker build -t ${algorithm_name} .`

# Bring Your Own Container - Steps

Code



Containerize



Register with  
Amazon ECR



Create  
Estimator

# Bring Your Own Container - Steps

Code



Containerize



Register with  
Amazon ECR




Create  
Estimator


- `aws ecr create-repository --repository-name "${algorithm_name}" > /dev/null`
- `fullname="${account}.dkr.ecr.${region}.amazonaws.com/${algorithm_name}:latest"`
- `docker push ${fullname}`

# Bring Your Own Container - Steps

Code



Containerize



Register with  
Amazon ECR



Create  
Estimator

# Bring Your Own Container - Steps

Code



Containerize



Register with  
Amazon ECR



Create  
Estimator

- `byoc_image_uri = '{}.dkr.ecr.{}.{}' / {}'.format(account_id, region, uri_suffix, ecr_repository + tag)`
- `estimator = Estimator ( image_name=byoc_image_uri, .... )`

# Summary





# Summary

- Tune and evaluate a model
- Model tuning
- Tune a BERT-based text classifier
- Model evaluation
- Evaluate a BERT-based text classifier
- TODO: Script Mode?
- TODO: Bring Your Own Container?

